

shap-and-lime-for-models

May 15, 2024

```
[ ]: !pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.35.2)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.13.1)
Requirement already satisfied: huggingface-hub<1.0,>=0.16.4 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.19.4)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.23.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (23.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.1)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2023.6.3)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.31.0)
Requirement already satisfied: tokenizers<0.19,>=0.14 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.15.0)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.4.1)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.66.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.16.4->transformers) (2023.6.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.16.4->transformers) (4.5.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->transformers)
```

(2023.11.17)

```
[ ]: import numpy as np
import re
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from transformers import BertTokenizer, TFBertForSequenceClassification
from transformers import TFBertForSequenceClassification
from sklearn.metrics import confusion_matrix, classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Bidirectional, GRU, Dropout, Dense
from tensorflow.keras.layers import Embedding
%matplotlib inline
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: df = pd.read_excel('/content/drive/MyDrive/Datasets/Bengali_comment.xlsx')
df.head()
```

```
[ ]:
```

				comment	Category	Gender	\
0		...	Actor	Female			
1	?	...	Singer	Male			
2		,	????	Actor	Female		
3				Sports	Male		
4				Politician	Male		

	comment	react	number	label
0		1.0		sexual
1		2.0		not bully
2		2.0		not bully
3		0.0		not bully
4		0.0		troll

```
[ ]: df.isnull().sum()
```

```
[ ]: comment          0
Category            0
Gender              0
comment react number 3
label               0
```


dtype: int64

```
[ ]: df.dropna(inplace=True)
```


```
[ ]: df['label'].value_counts()
```

```
[ ]: not bully    15339
      troll      10462
      sexual     8928
      religious   7575
      threat     1694
      Name: label, dtype: int64
```

```
[ ]: plt.figure(figsize=(10,4))
      sns.countplot(x='label',data=df)
      plt.xlabel('Label')
      plt.ylabel('Count')
      plt.title('Number of ham and spam messages')
      plt.show()
```




```
[ ]: plt.figure(figsize=(12, 6))
      sns.countplot(x='Category', data=df)
      plt.title('Distribution of Comments Across Categories')
      plt.xlabel('Category')
      plt.ylabel('Count')
      plt.show()
```




```
[ ]: # Step 2: Pie Charts for percentage distribution of labels within each category
plt.figure(figsize=(20, 15)) # Increase the figure size
categories = df['Category'].unique()
num_categories = len(categories)


for i, category in enumerate(categories):
    plt.subplot((num_categories // 2) + 1, 2, i+1)
    category_df = df[df['Category'] == category]
    labels = category_df['label'].unique()
    pie = plt.pie(category_df['label'].value_counts(), labels=labels,
    ↪ autopct='%1.1f%%', startangle=90, pctdistance=0.85)
    plt.title(f'Label Distribution in {category}')
    plt.legend(labels, loc="center left", bbox_to_anchor=(1, 0, 0.5, 1))
plt.tight_layout()
plt.show()
```



```
[ ]: plt.figure(figsize=(12, 6))
sns.boxplot(x='Category', y='comment react number', data=df)
plt.title('Distribution of React Numbers Within Each Category')
plt.xlabel('Category')
plt.ylabel('React Number')
plt.show()
```



```
[ ]: plt.figure(figsize=(12, 6))
sns.violinplot(x='Category', y='comment react number', data=df)
plt.title('Distribution of React Numbers Within Each Category (Violin Plot)')
plt.xlabel('Category')
plt.ylabel('React Number')
plt.show()
```



```
import nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
```

```
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

```
[ ]: True
```

```
import nltk
from nltk.corpus import stopwords
nltk.download('punkt')
nltk.download('stopwords')

bengali_stopwords = set(stopwords.words('bengali'))
print(bengali_stopwords)
df['comment'] = df['comment'].apply(lambda x: ' '.join([word for word in x.
    ↪split() if word not in bengali_stopwords]))
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
```

```
[nltk_data] Package punkt is already up-to-date!
```


```
[nltk_data] Downloading package stopwords to /root/nltk_data...
```

```
[nltk_data]   Unzipping corpora/stopwords.zip.
```


A large grid of 100 rows and 100 columns of curly braces '{' and '}' characters, creating a dense, textured pattern.

[illegible]

```
[ ]: # pie chart of the labels not bully troll sexual religious threat
plt.figure(figsize=(10, 10))
df['label'].value_counts().plot.pie(autopct='%1.1f%%')
plt.show()
```

```
[ ]: # text sequence length
plt.figure(figsize=(10,4))
df['length'] = df['comment'].apply(len)
sns.histplot(df['length'],kde=True,bins=100)
plt.xlabel('Length')
plt.ylabel('Count')
plt.title('Length of text sequence')
plt.show()
```



```
[ ]: # explore the datasets
def explore_data(data):
    for i in range(5):
        print("Sample Comment:-\n",data['comment'][i])
        print("-----")
        print("Sample Label:-\n",data['label'][i])
        print("-----")

    # analyse the length of text
    text_len = [len(text) for text in data['comment']]
    print("Average length of text:-",np.mean(text_len))
    print("Max length of text:-",np.max(text_len))
    print("Min length of text:-",np.min(text_len))
    print("Standard deviation of length of text:-",np.std(text_len))
    print("Median length of text:-",np.median(text_len))
    print("25 percentile of length of text:-",np.percentile(text_len,25))
    print("75 percentile of length of text:-",np.percentile(text_len,75))
    print("-----")
```

```
[ ]: explore_data(df)
```

Sample Comment:-

**** safa

Sample Label:-

sexual

Sample Comment:-

? ?


```
[ ]: for i in range(len(df)):
    text = df.loc[i, 'comment']
    for punctuation in remove_punctuations:
        text = text.replace(punctuation, ' ')
    df.loc[i, 'comment'] = text
```

```
[ ]: # remove emoji
def remove_emoji(text):
    emoji_pattern = re.compile(
        "[u"\U0001F600-\U0001F64F" # emoticons
        u"\U0001F300-\U0001F5FF" # symbols & pictographs
        u"\U0001F680-\U0001F6FF" # transport & map symbols
        u"\U0001F1E0-\U0001F1FF" # flags (iOS)
        u"\U00002702-\U000027B0"
        u"\U000024C2-\U0001F251"
        ]+",
        flags=re.UNICODE,
    )
    return emoji_pattern.sub(r"", text)
```

```
[ ]: # remove emoji
for i in range(len(df)):
    text = df.loc[i, 'comment']
    text = remove_emoji(text)
    df.loc[i, 'comment'] = text
```

```
[ ]: # remove english character
def remove_english_character(text):
    english_character = re.compile("[a-zA-Z]+")
    return english_character.sub(r"", text)
```

```
[ ]: # remove english character
for i in range(len(df)):
    text = df.loc[i, 'comment']
    text = remove_english_character(text)
    df.loc[i, 'comment'] = text
```

```
[ ]: # remove extra space
def remove_extra_space(text):
    extra_space = re.compile("\s+")
    return extra_space.sub(r" ", text)
```

```
[ ]: def remove_single_bengali_character(text):
    # Regular expression pattern to match single Bengali characters
    single_character = re.compile(r'\s[ -]\s')
    return single_character.sub(" ", text)
```

```
# Identify data to check if the remove_single_bengali_character function works
for i in range(5):
    print("Original data:-\n", df['comment'][i])
    print("Processed data:-\n",
    ↪remove_single_bengali_character(df['comment'][i]))
    print("-----")
```

Original data:-

Processed data:-

Original data:-

Processed data:-

Original data:-

Processed data:-

Original data:-

Processed data:-

Original data:-

Processed data:-

```
[ ]: df['comment'] = df['comment'].apply(remove_single_bengali_character)
df.head()
```

```
[ ]:
  index      comment  Category \
0      0      Actor
1      1      Singer
2      2      Actor
3      3      Sports
4      4      Politician

  Gender  comment  react  number  label  length
0  Female      1.0    sexual    140
```

1	Male	2.0	not bully	38
2	Female	2.0	not bully	21
3	Male	0.0	not bully	21
4	Male	0.0	troll	8

```
[ ]: explore_data(df)
```

Sample Comment:-

Sample Label:-
sexual

Sample Comment:-

Sample Label:-
not bully

Sample Comment:-

Sample Label:-
not bully

Sample Comment:-

Sample Label:-
not bully

Sample Comment:-

Sample Label:-
troll

Average length of text:- 74.27898995408883
 Max length of text:- 1319
 Min length of text:- 0
 Standard deviation of length of text:- 107.8391821371824
 Median length of text:- 39.0
 25 percentile of length of text:- 20.0
 75 percentile of length of text:- 82.0

```
[ ]: # remove extra space
for i in range(len(df)):
    text = df.loc[i, 'comment']
    text = remove_extra_space(text)
    df.loc[i, 'comment'] = text
```

```
[ ]: explore_data(df)
```

Sample Comment:-

Sample Label:-

sexual

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-

troll

Average length of text:- 71.9247011227783

Max length of text:- 1195


Min length of text:- 0


Standard deviation of length of text:- 103.40000622663459

Median length of text:- 38.0

25 percentile of length of text:- 20.0

75 percentile of length of text:- 79.0






```
print(len(unique_words))
```

56199

```
[ ]: # total number of words
total_words = [word for comment in df['comment'] for word in comment.split()]
print(len(total_words))
```

535585

```
[ ]: df = df[['comment', 'label']]
```

```
[ ]: df.head()
```

```
[ ]:
      ...      comment      label
0      sexual
1      not bully
2      not bully
3      not bully
4      troll
```

```
[ ]: explore_data(df)
```

Sample Comment:-

Sample Label:-

sexual

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-

not bully

Sample Comment:-

Sample Label:-
troll

Average length of text:- 71.9247011227783
Max length of text:- 1195
Min length of text:- 0
Standard deviation of length of text:- 103.40000622663459
Median length of text:- 38.0
25 percentile of length of text:- 20.0
75 percentile of length of text:- 79.0

```
[ ]: le = LabelEncoder()  
df['label'] = le.fit_transform(df['label'])  
  
labels = to_categorical(df['label'], num_classes=5)  
  
df.head()
```

```
[ ]:  
      ...      2      0      0      4  
0  
1  
2  
3      0  
4
```

```
[ ]: train_texts, test_texts, train_labels, test_labels =  
      ↪train_test_split(df['comment'].tolist(), df['label'].tolist(), test_size=0.2)
```

```
[ ]: from transformers import TFBertModel  
import tensorflow as tf
```

```
[ ]: tokenizer = BertTokenizer.from_pretrained("sagorsarker/bangla-bert-base")
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:72:  
UserWarning:  
The secret `HF_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab  
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab  
and restart your session.  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access  
public models or datasets.  
warnings.warn(  

```

```
vocab.txt: 0%|          | 0.00/2.24M [00:00<?, ?B/s]
config.json: 0%|          | 0.00/491 [00:00<?, ?B/s]
```

```
[ ]: max_length = 128
train_encodings = tokenizer(train_texts, truncation=True, padding=True,
    ↳max_length=max_length, return_tensors="tf")
test_encodings = tokenizer(test_texts, truncation=True, padding=True,
    ↳max_length=max_length, return_tensors="tf")
```

```
[ ]: num_labels = len(df['label'].unique())
```

```
[ ]: # Assuming train_encodings and test_encodings contain input_ids
train_input_ids = train_encodings['input_ids']
test_input_ids = test_encodings['input_ids']

# Trim or pad sequences to the desired length (128)
max_length = 128
train_input_ids = tf.keras.preprocessing.sequence.
    ↳pad_sequences(train_input_ids, maxlen=max_length, padding='post')
test_input_ids = tf.keras.preprocessing.sequence.pad_sequences(test_input_ids,
    ↳maxlen=max_length, padding='post')
```

```
[ ]: # Create datasets
train_dataset = tf.data.Dataset.from_tensor_slices((
    {
        'sequences': train_input_ids,
        'attention_mask': train_encodings['attention_mask']
    },
    tf.keras.utils.to_categorical(train_labels, num_labels)
))

test_dataset = tf.data.Dataset.from_tensor_slices((
    {
        'sequences': test_input_ids,
        'attention_mask': test_encodings['attention_mask']
    },
    tf.keras.utils.to_categorical(test_labels, num_labels)
))
```

```
[ ]: # One-hot encode labels
train_labels_onehot = to_categorical(train_labels, num_labels)
test_labels_onehot = to_categorical(test_labels, num_labels)
```

1 GRU Model

```
[ ]: from keras.layers import Input, Embedding, Bidirectional, GRU, Dense, Dropout, \
    ↳ Attention, Reshape
from keras.models import Model
from keras.layers import GlobalAveragePooling1D
from keras.layers import Input, Embedding, Bidirectional, GRU, Dense, Dropout, \
    ↳ Attention, Reshape, Flatten
from keras.models import Model
```

```
[ ]: vocab_size = tokenizer.vocab_size
embedding_dim = 300
```

```
[ ]: def GRUmodel(vocab_size, embedding_dim=128, sequence_length=128):
    sequences = Input(shape=(sequence_length,), dtype='int32', name='sequences')

    embedded_sequences = Embedding(vocab_size, embedding_dim, \
    ↳ input_length=sequence_length)(sequences)

    # Enhanced GRU layers
    x = Bidirectional(GRU(256, return_sequences=True))(embedded_sequences)
    x = Bidirectional(GRU(128))(x)
    x = Dense(64, activation='relu')(x)
    x = Dropout(0.2)(x)
    # Attention mechanism
    attention = Attention()([x, x])

    # Flatten and additional dense layers
    x = Flatten()(attention)
    x = Dense(128, activation='relu')(x)
    x = Dropout(0.2)(x)
    x = Dense(64, activation='relu')(x)
    x = Dropout(0.1)(x)
    x = Dense(32, activation='relu')(x)

    num_classes = 5
    output = Dense(num_classes, activation='softmax')(x)

    return Model(inputs=sequences, outputs=output)
```

```
[ ]: vocab_size = tokenizer.vocab_size
model = GRUmodel(vocab_size, sequence_length=128)
model.summary()
```

Model: "model_3"

Layer (type)	Output Shape	Param #	Connected to
=====			
sequences (InputLayer)	[(None, 128)]	0	[]
embedding_3 (Embedding)	(None, 128, 128)	1305280	
['sequences[0][0]']		0	
bidirectional_5 (Bidirectional)	(None, 128, 512)	592896	
['embedding_3[0][0]']			
bidirectional_6 (Bidirectional)	(None, 256)	493056	
['bidirectional_5[0][0]']			
dense_10 (Dense)	(None, 64)	16448	
['bidirectional_6[0][0]']			
dropout_5 (Dropout)	(None, 64)	0	
['dense_10[0][0]']			
attention_3 (Attention)	(None, 64)	0	
['dropout_5[0][0]', 'dropout_5[0][0]']			
flatten_3 (Flatten)	(None, 64)	0	
['attention_3[0][0]']			
dense_11 (Dense)	(None, 128)	8320	
['flatten_3[0][0]']			
dropout_6 (Dropout)	(None, 128)	0	
['dense_11[0][0]']			
dense_12 (Dense)	(None, 64)	8256	
['dropout_6[0][0]']			
dropout_7 (Dropout)	(None, 64)	0	
['dense_12[0][0]']			
dense_13 (Dense)	(None, 32)	2080	
['dropout_7[0][0]']			
dense_14 (Dense)	(None, 5)	165	
['dense_13[0][0]']			

```
=====
=====
Total params: 14174021 (54.07 MB)
Trainable params: 14174021 (54.07 MB)
Non-trainable params: 0 (0.00 Byte)
-----
-----
```

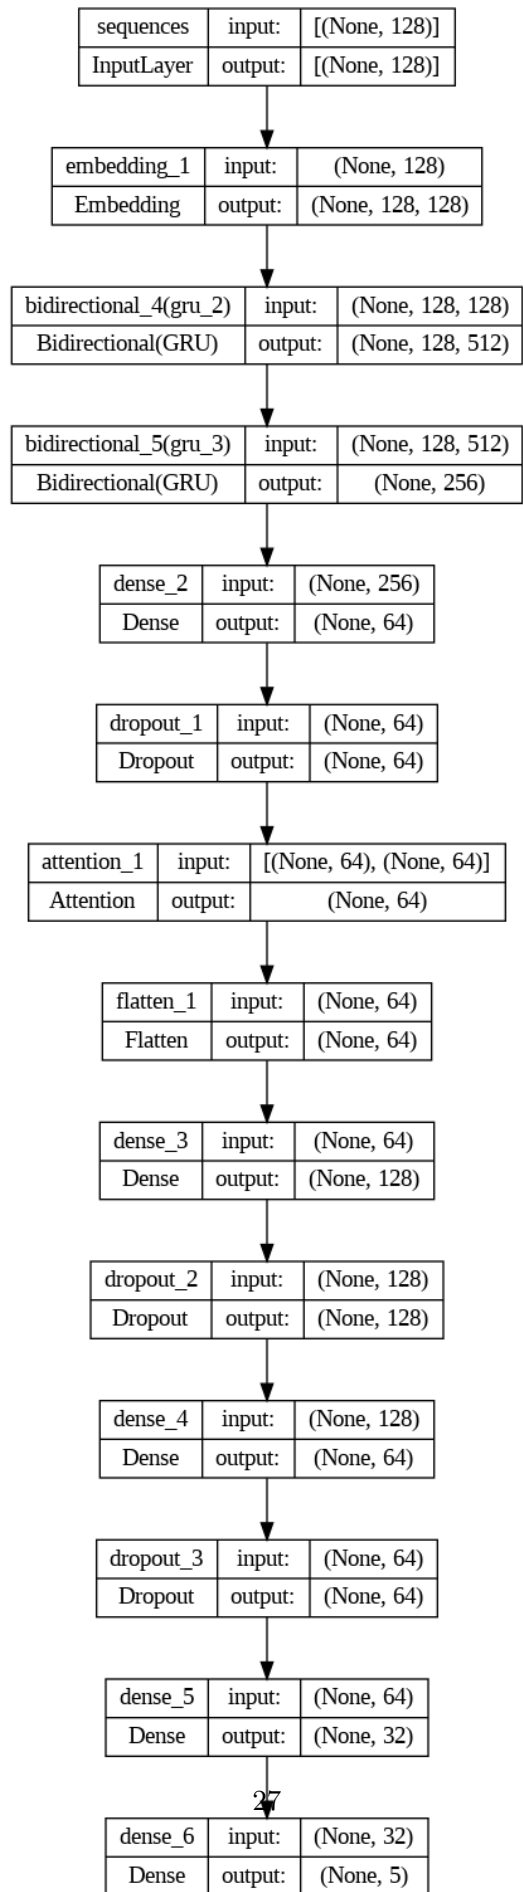
```
[ ]: !pip install pydot graphviz
```

```
Requirement already satisfied: pydot in /usr/local/lib/python3.10/dist-packages
(1.4.2)
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-
packages (0.20.1)
Requirement already satisfied: pyparsing>=2.1.4 in
/usr/local/lib/python3.10/dist-packages (from pydot) (3.1.1)
```

```
[ ]: from tensorflow.keras.utils import plot_model

# Save the model summary as an image file
plot_model(model, to_file='hybrid_model_summary.png', show_shapes=True,
↪show_layer_names=True)
```

```
[ ]:
```



```
[ ]: from keras.optimizers import Adam

optimizer = Adam(learning_rate=0.001) # Adjust the learning rate if needed
model.compile(optimizer=optimizer, loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[ ]: train_labels_onehot = tf.keras.utils.to_categorical(train_labels,
    ↪num_classes=len(set(train_labels)))
test_labels_onehot = tf.keras.utils.to_categorical(test_labels,
    ↪num_classes=len(set(test_labels)))
```

```
[ ]: from tensorflow.keras.callbacks import EarlyStopping
```

```
[ ]: # Early Stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=3,
    ↪restore_best_weights=True)
```

```
[ ]: # Train the model
history = model.fit(
    train_encodings['input_ids'],
    train_labels_onehot,
    validation_data=(test_encodings['input_ids'],
        test_labels_onehot),
    epochs=10,
    batch_size=32,
    callbacks=[early_stopping]
)
```

```
Epoch 1/10
1100/1100 [=====] - 192s 175ms/step - loss: 0.6115 -
accuracy: 0.8265 - val_loss: 0.7033 - val_accuracy: 0.7886
Epoch 2/10
1100/1100 [=====] - 192s 174ms/step - loss: 0.5562 -
accuracy: 0.8418 - val_loss: 0.6532 - val_accuracy: 0.7965
Epoch 3/10
1100/1100 [=====] - 193s 175ms/step - loss: 0.5390 -
accuracy: 0.8481 - val_loss: 0.6810 - val_accuracy: 0.7973
Epoch 4/10
1100/1100 [=====] - 184s 168ms/step - loss: 0.5042 -
accuracy: 0.8592 - val_loss: 0.6760 - val_accuracy: 0.7969
Epoch 5/10
1100/1100 [=====] - 183s 167ms/step - loss: 0.4872 -
accuracy: 0.8646 - val_loss: 0.6531 - val_accuracy: 0.7994
Epoch 6/10
1100/1100 [=====] - 186s 169ms/step - loss: 0.4632 -
```

```

accuracy: 0.8703 - val_loss: 0.6715 - val_accuracy: 0.7973
Epoch 7/10
1100/1100 [=====] - 192s 175ms/step - loss: 0.4593 -
accuracy: 0.8715 - val_loss: 0.6453 - val_accuracy: 0.7995
Epoch 8/10
1100/1100 [=====] - 192s 175ms/step - loss: 0.4601 -
accuracy: 0.8752 - val_loss: 0.6489 - val_accuracy: 0.8070
Epoch 9/10
1100/1100 [=====] - 184s 167ms/step - loss: 0.4187 -
accuracy: 0.8863 - val_loss: 0.6570 - val_accuracy: 0.8073
Epoch 10/10
1100/1100 [=====] - 184s 168ms/step - loss: 0.3990 -
accuracy: 0.8916 - val_loss: 0.6624 - val_accuracy: 0.8031

```

```

[ ]: # Evaluate the model on test data
loss, accuracy = model.evaluate(test_dataset.batch(32))
print(f'Test Accuracy: {accuracy * 100:.2f}%')
print(f'Test Loss: {loss:.4f}')

# Get predictions for test data
predictions = model.predict(test_dataset.batch(32))
predicted_classes = predictions.argmax(axis=1)
true_classes = test_labels

```

```

4/275 [...] - ETA: 13s - loss: 0.8858 - accuracy:
0.7188

```

```

/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
UserWarning: Input dict contained keys ['attention_mask'] which did not match
any model input. They will be ignored by the model.

```

```

inputs = self._flatten_to_reference_inputs(inputs)

275/275 [=====] - 14s 51ms/step - loss: 0.6453 -
accuracy: 0.7995
Test Accuracy: 79.95%
Test Loss: 0.6453
275/275 [=====] - 10s 36ms/step

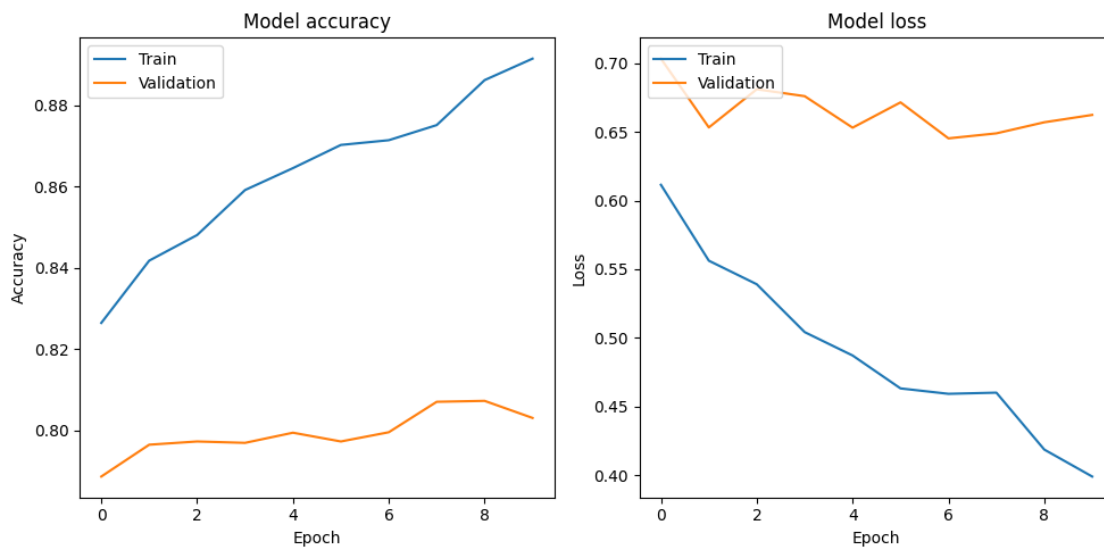
```

```

[ ]: # Plot training & validation accuracy values
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

```

```
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.tight_layout()
plt.show()
```



```
[ ]: from sklearn.metrics import classification_report, confusion_matrix,
precision_recall_fscore_support
```

```
[ ]: # Classification Report
class_report = classification_report(true_classes, predicted_classes)
print("Classification Report:\n", class_report)

# Confusion Matrix
conf_matrix = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:\n", conf_matrix)

# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
```

```
plt.show()


# Precision, Recall, F1-score
precision, recall, f1_score, _ = precision_recall_fscore_support(true_classes,
    ↪ predicted_classes, average='weighted')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1_score:.4f}')
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.82	0.83	3089
1	0.92	0.85	0.88	1528
2	0.80	0.82	0.81	1746
3	0.83	0.49	0.61	349
4	0.67	0.76	0.72	2088
accuracy			0.80	8800
macro avg	0.81	0.75	0.77	8800
weighted avg	0.81	0.80	0.80	8800

Confusion Matrix:

```
[[2547  14  95   3 430]
 [ 62 1296  75   7  88]
 [ 102  19 1427   3 195]
 [ 40  55  19 170  65]
 [ 283  20 168  21 1596]]
```



Precision: 0.8062

Recall: 0.7995

F1-score: 0.8001

```
[ ]: # Extract TP, TN, FP, FN from confusion matrix
TP = conf_matrix[1, 1] # True Positives
TN = conf_matrix[0, 0] # True Negatives
FP = conf_matrix[0, 1] # False Positives
FN = conf_matrix[1, 0] # False Negatives

print("True Positives:", TP)
print("True Negatives:", TN)
print("False Positives:", FP)
print("False Negatives:", FN)
```

True Positives: 1296

True Negatives: 2547

False Positives: 14

False Negatives: 62


```
[ ]: from sklearn.preprocessing import LabelBinarizer
      from sklearn.metrics import roc_curve, roc_auc_score

      # Get predictions for test data
      predictions = model.predict(test_dataset.batch(32))
      predicted_classes = predictions.argmax(axis=1)
      true_classes = test_labels

      # Transform true labels to binary format
      label_binarizer = LabelBinarizer()
      true_labels_bin = label_binarizer.fit_transform(true_classes)

      # Calculate ROC curve and AUC for each class
      fpr = dict()
      tpr = dict()
      roc_auc = dict()
      num_labels = len(label_binarizer.classes_)

      for i in range(num_labels):
          fpr[i], tpr[i], _ = roc_curve(true_labels_bin[:, i], predictions[:, i])
          roc_auc[i] = roc_auc_score(true_labels_bin[:, i], predictions[:, i])

      # Plot ROC curve for each class
      plt.figure(figsize=(8, 6))

      for i in range(num_labels):
          plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')


      plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve for Each Class')
      plt.legend(loc="lower right")
      plt.show()
```

3/275 [...] - ETA: 14s

/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
 UserWarning: Input dict contained keys ['attention_mask'] which did not match
 any model input. They will be ignored by the model.

inputs = self._flatten_to_reference_inputs(inputs)

275/275 [=====] - 11s 38ms/step



```
[ ]: from keras.layers import Input, Embedding, Bidirectional, GRU, LSTM, Dense,
      Flatten, Dropout, Attention, Concatenate
from keras.models import Model
from keras.optimizers import Adam

def HybridModel(vocab_size, embedding_dim=128, sequence_length=128):
    sequences = Input(shape=(sequence_length,), dtype='int32', name='sequences')

    embedded_sequences = Embedding(vocab_size, embedding_dim,
    input_length=sequence_length)(sequences)

    # GRU layers
    gru_out = Bidirectional(GRU(128, return_sequences=True))(embedded_sequences)
    gru_out = Bidirectional(GRU(64))(gru_out)

    # LSTM layers
    lstm_out = Bidirectional(LSTM(128,
    return_sequences=True))(embedded_sequences)
    lstm_out = Bidirectional(LSTM(64))(lstm_out)
```

```

# Concatenate GRU and LSTM outputs
concatenated = Concatenate()([gru_out, lstm_out])

# Attention mechanism
attention = Attention()([concatenated, concatenated])

# Flatten and additional dense layers
x = Flatten()(attention)
x = Dense(128, activation='relu')(x)
x = Dropout(0.5)(x)
x = Dense(64, activation='relu')(x)
x = Dropout(0.3)(x)
x = Dense(32, activation='relu')(x)

num_classes = 5
output = Dense(num_classes, activation='softmax')(x)

model = Model(inputs=sequences, outputs=output)
optimizer = Adam(learning_rate=0.001) # Adjust the learning rate if needed
model.compile(optimizer=optimizer, loss='categorical_crossentropy',
metrics=['accuracy'])

return model

```

```

[ ]: vocab_size = tokenizer.vocab_size
model = HybridModel(vocab_size, sequence_length=128)
model.summary()

```

Model: "model"

```

-----
Layer (type)                Output Shape              Param #   Connected to
=====
sequences (InputLayer)      [(None, 128)]             0         []
embedding (Embedding)       (None, 128, 128)          1305280   ['sequences[0][0]']
                                0
bidirectional (Bidirection  (None, 128, 256)          198144    ['embedding[0][0]']
al)
bidirectional_2 (Bidirecti  (None, 128, 256)          263168    ['embedding[0][0]']

```

```

onal)

bidirectional_1 (Bidirecti (None, 128) 123648
['bidirectional[0][0]']
onal)

bidirectional_3 (Bidirecti (None, 128) 164352
['bidirectional_2[0][0]']
onal)

concatenate (Concatenate) (None, 256) 0
['bidirectional_1[0][0]',
'bidirectional_3[0][0]']

attention (Attention) (None, 256) 0
['concatenate[0][0]',
'concatenate[0][0]']

flatten (Flatten) (None, 256) 0
['attention[0][0]']

dense (Dense) (None, 128) 32896
['flatten[0][0]']

dropout (Dropout) (None, 128) 0
['dense[0][0]']

dense_1 (Dense) (None, 64) 8256
['dropout[0][0]']

dropout_1 (Dropout) (None, 64) 0
['dense_1[0][0]']

dense_2 (Dense) (None, 32) 2080
['dropout_1[0][0]']

dense_3 (Dense) (None, 5) 165
['dense_2[0][0]']

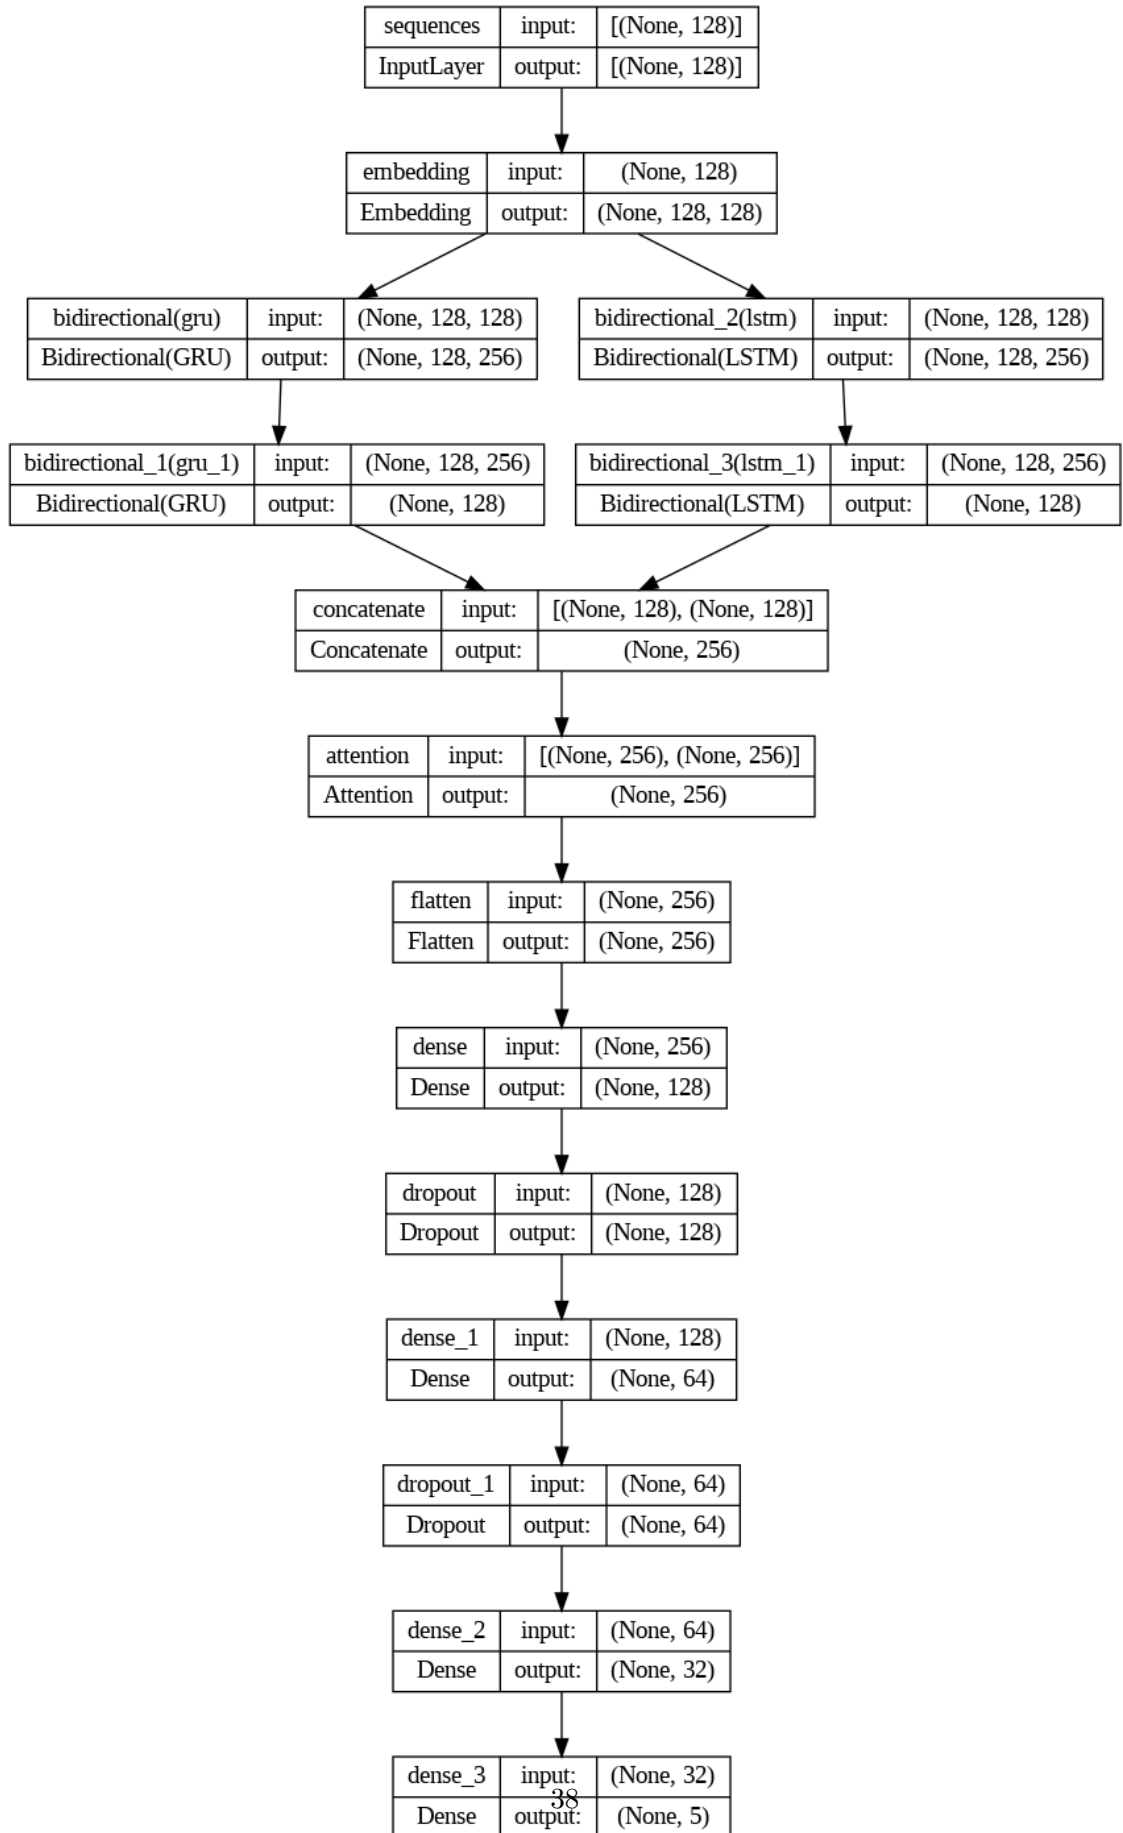
=====
=====
Total params: 13845509 (52.82 MB)
Trainable params: 13845509 (52.82 MB)
Non-trainable params: 0 (0.00 Byte)
-----
-----

```

```
[ ]: from tensorflow.keras.utils import plot_model

# Save the model summary as an image file
plot_model(model, to_file='hybrid_model_summary.png', show_shapes=True,
            show_layer_names=True)
```

```
[ ]:
```



```
[ ]: # Early Stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
    ↪restore_best_weights=True)
```

```
[ ]: # Train the model
history = model.fit(
    train_encodings['input_ids'],
    train_labels_onehot,
    validation_data=(test_encodings['input_ids'],
                    test_labels_onehot),
    epochs=20,
    batch_size=32,
    callbacks=[early_stopping]
)
```

```
Epoch 1/20
1100/1100 [=====] - 100s 76ms/step - loss: 1.2824 -
accuracy: 0.4571 - val_loss: 0.8415 - val_accuracy: 0.7123
Epoch 2/20
1100/1100 [=====] - 56s 51ms/step - loss: 0.7167 -
accuracy: 0.7710 - val_loss: 0.6251 - val_accuracy: 0.7972
Epoch 3/20
1100/1100 [=====] - 51s 46ms/step - loss: 0.5304 -
accuracy: 0.8363 - val_loss: 0.5625 - val_accuracy: 0.8214
Epoch 4/20
1100/1100 [=====] - 52s 47ms/step - loss: 0.4316 -
accuracy: 0.8671 - val_loss: 0.5731 - val_accuracy: 0.8190
Epoch 5/20
1100/1100 [=====] - 50s 46ms/step - loss: 0.3559 -
accuracy: 0.8894 - val_loss: 0.5697 - val_accuracy: 0.8224
Epoch 6/20
1100/1100 [=====] - 52s 47ms/step - loss: 0.2980 -
accuracy: 0.9085 - val_loss: 0.6063 - val_accuracy: 0.8184
Epoch 7/20
1100/1100 [=====] - 52s 47ms/step - loss: 0.2504 -
accuracy: 0.9221 - val_loss: 0.6499 - val_accuracy: 0.8193
Epoch 8/20
1100/1100 [=====] - 51s 46ms/step - loss: 0.2105 -
accuracy: 0.9354 - val_loss: 0.7081 - val_accuracy: 0.8120
```

```
[ ]: # Save the entire model to a HDF5 file
model.save('/content/drive/MyDrive/Bully/GRU_Model')
```

```
[ ]: # Evaluate the model on test data
loss, accuracy = model.evaluate(test_dataset.batch(32))
print(f'Test Accuracy: {accuracy * 100:.2f}%')
print(f'Test Loss: {loss:.4f}')

# Get predictions for test data
predictions = model.predict(test_dataset.batch(32))
predicted_classes = predictions.argmax(axis=1)
true_classes = test_labels
```

/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
UserWarning: Input dict contained keys ['attention_mask'] which did not match
any model input. They will be ignored by the model.

```
inputs = self._flatten_to_reference_inputs(inputs)
```

275/275 [=====] - 7s 17ms/step - loss: 0.5625 -
accuracy: 0.8214


Test Accuracy: 82.14%

Test Loss: 0.5625

275/275 [=====] - 7s 15ms/step

```
[ ]: # Plot training & validation accuracy values
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.tight_layout()
plt.show()
```

```
[ ]: from sklearn.metrics import classification_report, confusion_matrix, \
      ↪ precision_recall_fscore_support
```


```
[ ]: # Classification Report
class_report = classification_report(true_classes, predicted_classes)
print("Classification Report:\n", class_report)
# Confusion Matrix
conf_matrix = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:\n", conf_matrix)
# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.92	0.84	3129
1	0.95	0.84	0.89	1466
2	0.84	0.80	0.82	1784
3	0.83	0.65	0.73	357
4	0.81	0.71	0.75	2064
accuracy			0.82	8800
macro avg	0.84	0.78	0.81	8800
weighted avg	0.83	0.82	0.82	8800

Confusion Matrix:

```
[[2871  18  75   6 159]
 [ 130 1238  42  21  35]
 [ 210  16 1432   5 121]
 [  61  17  11 231  37]
 [ 428  12  153  15 1456]]
```



```
[ ]: # Extract TP, TN, FP, FN from confusion matrix
TP = conf_matrix[1, 1] # True Positives
TN = conf_matrix[0, 0] # True Negatives
FP = conf_matrix[0, 1] # False Positives
FN = conf_matrix[1, 0] # False Negatives

print("True Positives:", TP)
print("True Negatives:", TN)
print("False Positives:", FP)
print("False Negatives:", FN)
```

True Positives: 1238

True Negatives: 2871
False Positives: 18
False Negatives: 130

```
[ ]: # Precision, Recall, F1-score
precision, recall, f1_score, _ = precision_recall_fscore_support(true_classes,
    predicted_classes, average='weighted')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1_score:.4f}')
```

Precision: 0.8265
Recall: 0.8214
F1-score: 0.8200

```
[ ]: from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import roc_curve, roc_auc_score
```

```
[ ]: from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import roc_curve, roc_auc_score

# Get predictions for test data
predictions = model.predict(test_dataset.batch(32))
predicted_classes = predictions.argmax(axis=1)
true_classes = test_labels

# Transform true labels to binary format
label_binarizer = LabelBinarizer()
true_labels_bin = label_binarizer.fit_transform(true_classes)

# Calculate ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
num_labels = len(label_binarizer.classes_)

for i in range(num_labels):
    fpr[i], tpr[i], _ = roc_curve(true_labels_bin[:, i], predictions[:, i])
    roc_auc[i] = roc_auc_score(true_labels_bin[:, i], predictions[:, i])


# Plot ROC curve for each class
plt.figure(figsize=(8, 6))

for i in range(num_labels):
    plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Each Class')
plt.legend(loc="lower right")
plt.show()
```

275/275 [=====] - 5s 18ms/step



LSTM Model

```
[ ]: from tensorflow.keras.layers import Input, Embedding, LSTM, Dropout, Dense,
↳ Bidirectional
from tensorflow.keras.models import Model
```

```
[ ]: def LSTMmodel(vocab_size, embedding_dim=128, sequence_length=128):
    # Define input layer
    sequences = Input(shape=(sequence_length,), dtype=tf.int32,
↳ name="sequences")

    embedded_sequences = Embedding(vocab_size, embedding_dim)(sequences)
```

```

# LSTM layers
lstm_out = Bidirectional(LSTM(128,
↪return_sequences=True))(embedded_sequences)
lstm_out = Dropout(0.5)(lstm_out)
lstm_out = LSTM(64, return_sequences=True)(lstm_out)
lstm_out = Dropout(0.5)(lstm_out)

# Attention layer
attention = Attention()([lstm_out, lstm_out])
x = Flatten()(attention)

# Dense layers
x = Dense(64, activation='relu')(x)
x = Dropout(0.5)(x)
num_classes = len(set(train_labels))
x = Dense(num_classes, activation='softmax')(x)

return Model(inputs=sequences, outputs=x)

```

```

[ ]: # Create the model
vocab_size = tokenizer.vocab_size
model = LSTMmodel(vocab_size)
print(model.summary())

```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
sequences (InputLayer)	[(None, 128)]	0	[]
embedding_1 (Embedding)	(None, 128, 128)	1305280	['sequences[0][0]']
bidirectional_4 (Bidirectional)	(None, 128, 256)	263168	['embedding_1[0][0]']
dropout_2 (Dropout)	(None, 128, 256)	0	['bidirectional_4[0][0]']
lstm_3 (LSTM)	(None, 128, 64)	82176	['dropout_2[0][0]']

dropout_3 (Dropout) ['lstm_3[0][0]']	(None, 128, 64)	0
attention_1 (Attention) ['dropout_3[0][0]', 'dropout_3[0][0]']	(None, 128, 64)	0
flatten_1 (Flatten) ['attention_1[0][0]']	(None, 8192)	0
dense_4 (Dense) ['flatten_1[0][0]']	(None, 64)	524352
dropout_4 (Dropout) ['dense_4[0][0]']	(None, 64)	0
dense_5 (Dense) ['dropout_4[0][0]']	(None, 5)	325

```
=====
Total params: 13922821 (53.11 MB)
Trainable params: 13922821 (53.11 MB)
Non-trainable params: 0 (0.00 Byte)
-----
None
```

```
[ ]: model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[ ]: # Train the model
history_1 = model.fit(
    train_encodings['input_ids'],
    train_labels_onehot,
    validation_data=(test_encodings['input_ids'],
                     test_labels_onehot),
    epochs=10,
    batch_size=32,
    callbacks=[early_stopping]
)
```

```
Epoch 1/10
1100/1100 [=====] - 64s 51ms/step - loss: 0.9375 -
accuracy: 0.6508 - val_loss: 0.6198 - val_accuracy: 0.7951
Epoch 2/10
1100/1100 [=====] - 29s 26ms/step - loss: 0.5585 -
accuracy: 0.8245 - val_loss: 0.5652 - val_accuracy: 0.8172
```

```

Epoch 3/10
1100/1100 [=====] - 26s 24ms/step - loss: 0.4400 -
accuracy: 0.8616 - val_loss: 0.5959 - val_accuracy: 0.8090
Epoch 4/10
1100/1100 [=====] - 25s 23ms/step - loss: 0.3581 -
accuracy: 0.8868 - val_loss: 0.6519 - val_accuracy: 0.8116
Epoch 5/10
1100/1100 [=====] - 25s 23ms/step - loss: 0.3004 -
accuracy: 0.9024 - val_loss: 0.6754 - val_accuracy: 0.8073
Epoch 6/10
1100/1100 [=====] - 25s 23ms/step - loss: 0.2540 -
accuracy: 0.9172 - val_loss: 0.7550 - val_accuracy: 0.8067
Epoch 7/10
1100/1100 [=====] - 25s 23ms/step - loss: 0.2218 -
accuracy: 0.9279 - val_loss: 0.8951 - val_accuracy: 0.7985

```

```


[ ]: # Plotting training & validation accuracy
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history_1.history['accuracy'], label='Training Accuracy')
plt.plot(history_1.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')

# Plotting training & validation loss
plt.subplot(1, 2, 2)
plt.plot(history_1.history['loss'], label='Training Loss')
plt.plot(history_1.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss')

plt.tight_layout()
plt.show()

```



```
[ ]: # Classification Report
class_report = classification_report(true_classes, predicted_classes)
print("Classification Report:\n", class_report)
# Confusion Matrix
conf_matrix = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:\n", conf_matrix)
# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

Classification Report:


	precision	recall	f1-score	support
0	0.78	0.92	0.84	3129
1	0.95	0.84	0.89	1466
2	0.84	0.80	0.82	1784
3	0.83	0.65	0.73	357
4	0.81	0.71	0.75	2064
accuracy			0.82	8800
macro avg	0.84	0.78	0.81	8800
weighted avg	0.83	0.82	0.82	8800

Confusion Matrix:

```
[[2871  18  75   6 159]
 [ 130 1238  42  21  35]
 [ 210  16 1432   5 121]
```



```
[ 61  17  11 231  37]
[ 428 12 153  15 1456]]
```



```
[ ]: # Extract TP, TN, FP, FN from confusion matrix
TP = conf_matrix[1, 1] # True Positives
TN = conf_matrix[0, 0] # True Negatives
FP = conf_matrix[0, 1] # False Positives
FN = conf_matrix[1, 0] # False Negatives

print("True Positives:", TP)
print("True Negatives:", TN)
print("False Positives:", FP)
print("False Negatives:", FN)
```

```
True Positives: 1238
True Negatives: 2871
False Positives: 18
False Negatives: 130
```

```
[ ]: # Evaluate the model on test data
loss, accuracy = model.evaluate(test_dataset.batch(32))
print(f'Test Accuracy: {accuracy * 100:.2f}%')
print(f'Test Loss: {loss:.4f}')
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
UserWarning: Input dict contained keys ['attention_mask'] which did not match
any model input. They will be ignored by the model.
    inputs = self._flatten_to_reference_inputs(inputs)

275/275 [=====] - 3s 8ms/step - loss: 0.5652 -
accuracy: 0.8172
Test Accuracy: 81.72%
Test Loss: 0.5652
```

```
[ ]: # Precision, Recall, F1-score
precision, recall, f1_score, _ = precision_recall_fscore_support(true_classes,
    ↪ predicted_classes, average='weighted')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1_score:.4f}')
```

```
Precision: 0.8265
Recall: 0.8214
F1-score: 0.8200
```

```
[ ]: # Save the entire model to a HDF5 file
model.save('/content/drive/MyDrive/Bully/LSTM_Model')
```

```
[ ]: # Get predictions for test data
predictions = model.predict(test_dataset.batch(32))
predicted_classes = predictions.argmax(axis=1)
true_classes = test_labels

# Transform true labels to binary format
label_binarizer = LabelBinarizer()
true_labels_bin = label_binarizer.fit_transform(true_classes)

# Calculate ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
num_labels = len(label_binarizer.classes_)

for i in range(num_labels):
    fpr[i], tpr[i], _ = roc_curve(true_labels_bin[:, i], predictions[:, i])
    roc_auc[i] = roc_auc_score(true_labels_bin[:, i], predictions[:, i])
```

```

# Plot ROC curve for each class
plt.figure(figsize=(8, 6))

for i in range(num_labels):
    plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')


plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Each Class')
plt.legend(loc="lower right")
plt.show()

```

/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
UserWarning: Input dict contained keys ['attention_mask'] which did not match
any model input. They will be ignored by the model.

inputs = self._flatten_to_reference_inputs(inputs)

275/275 [=====] - 4s 9ms/step



CNN Model

```
[ ]: from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten

[ ]: def CNNmodel(vocab_size, embedding_dim=128, sequence_length=128):
    # Define input layer
    sequences = Input(shape=(sequence_length,), dtype=tf.int32,
        ↪name="sequences")

    embedded_sequences = Embedding(vocab_size, embedding_dim)(sequences)

    # 1D Convolution layers
    x = Conv1D(128, 5, activation='relu')(embedded_sequences)
    x = MaxPooling1D(5)(x)
    x = Conv1D(128, 5, activation='relu')(x)
    x = MaxPooling1D(5)(x)

    # Attention layer
    attention = Attention()([x, x])

    # GlobalMaxPooling1D
    x = Flatten()(attention)
    x = Dense(128, activation='relu')(x)
    x = Dense(64, activation='relu')(x)
    x = Dense(32, activation='relu')(x)

    # Output layer
    num_classes = len(set(train_labels))
    x = Dense(num_classes, activation='softmax')(x)

    return Model(inputs=sequences, outputs=x)

[ ]: # Create the model
vocab_size = tokenizer.vocab_size
model = CNNmodel(vocab_size)
print(model.summary())
```

Model: "model_2"

```
-----
Layer (type)                 Output Shape              Param #   Connected to
=====
sequences (InputLayer)       [(None, 128)]             0         []
embedding_2 (Embedding)      (None, 128, 128)          1305280   ['sequences[0][0]']
                                0
```

conv1d (Conv1D)	(None, 124, 128)	82048
['embedding_2[0][0]']		
max_pooling1d (MaxPooling1D)	(None, 24, 128)	0
['conv1d[0][0]']		
conv1d_1 (Conv1D)	(None, 20, 128)	82048
['max_pooling1d[0][0]']		
max_pooling1d_1 (MaxPooling1D)	(None, 4, 128)	0
['conv1d_1[0][0]']		
attention_2 (Attention)	(None, 4, 128)	0
['max_pooling1d_1[0][0]', 'max_pooling1d_1[0][0]']		
flatten_2 (Flatten)	(None, 512)	0
['attention_2[0][0]']		
dense_6 (Dense)	(None, 128)	65664
['flatten_2[0][0]']		
dense_7 (Dense)	(None, 64)	8256
['dense_6[0][0]']		
dense_8 (Dense)	(None, 32)	2080
['dense_7[0][0]']		
dense_9 (Dense)	(None, 5)	165
['dense_8[0][0]']		

```
=====
=====
Total params: 13293061 (50.71 MB)
Trainable params: 13293061 (50.71 MB)
Non-trainable params: 0 (0.00 Byte)

-----
-----
None
```

```
[ ]: model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[ ]: from tensorflow.keras.callbacks import EarlyStopping

[ ]: early_stopping = EarlyStopping(monitor='val_loss', patience=3,
    ↪restore_best_weights=True, verbose=1)

[ ]: history_2 = model.fit(train_encodings['input_ids'], train_labels_onehot,
    ↪validation_data=(test_encodings['input_ids'],
    ↪test_labels_onehot),
    ↪epochs=10, batch_size=32, callbacks=[early_stopping])
```

```
Epoch 1/10
1100/1100 [=====] - 63s 50ms/step - loss: 0.8079 -
accuracy: 0.6963 - val_loss: 0.5978 - val_accuracy: 0.8009
Epoch 2/10
1100/1100 [=====] - 14s 13ms/step - loss: 0.4800 -
accuracy: 0.8476 - val_loss: 0.5734 - val_accuracy: 0.8130
Epoch 3/10
1100/1100 [=====] - 13s 12ms/step - loss: 0.3430 -
accuracy: 0.8922 - val_loss: 0.6309 - val_accuracy: 0.8140
Epoch 4/10
1100/1100 [=====] - 12s 11ms/step - loss: 0.2409 -
accuracy: 0.9260 - val_loss: 0.7148 - val_accuracy: 0.8061
Epoch 5/10
1099/1100 [=====>.] - ETA: 0s - loss: 0.1774 - accuracy:
0.9455Restoring model weights from the end of the best epoch: 2.
1100/1100 [=====] - 11s 10ms/step - loss: 0.1773 -
accuracy: 0.9455 - val_loss: 0.8218 - val_accuracy: 0.7955
Epoch 5: early stopping
```


```
[ ]: plt.figure(figsize=(12, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history_2.history['accuracy'], label='Training Accuracy')
plt.plot(history_2.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history_2.history['loss'], label='Training Loss')
plt.plot(history_2.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
```

```
plt.title('Training and Validation Loss')

plt.tight_layout()
plt.show()
```



```
[ ]: # Get predictions for test data
predictions = model.predict(test_dataset.batch(32))
predicted_classes = predictions.argmax(axis=1)
true_classes = test_labels

# Transform true labels to binary format
label_binarizer = LabelBinarizer()
true_labels_bin = label_binarizer.fit_transform(true_classes)

# Calculate ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
num_labels = len(label_binarizer.classes_)

for i in range(num_labels):
    fpr[i], tpr[i], _ = roc_curve(true_labels_bin[:, i], predictions[:, i])
    roc_auc[i] = roc_auc_score(true_labels_bin[:, i], predictions[:, i])

# Plot ROC curve for each class
plt.figure(figsize=(8, 6))

for i in range(num_labels):
    plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
```


```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Each Class')
plt.legend(loc="lower right")
plt.show()
```

53/275 [====>...] - ETA: 0s

/usr/local/lib/python3.10/dist-packages/keras/src/engine/functional.py:642:
UserWarning: Input dict contained keys ['attention_mask'] which did not match
any model input. They will be ignored by the model.

```
inputs = self._flatten_to_reference_inputs(inputs)
```

275/275 [=====] - 1s 2ms/step



```
[ ]: # Evaluate the model on test data
loss, accuracy = model.evaluate(test_dataset.batch(32))
print(f'Test Accuracy: {accuracy * 100:.2f}%')
print(f'Test Loss: {loss:.4f}')
```



```

# Classification Report
class_report = classification_report(true_classes, predicted_classes)
print("Classification Report:\n", class_report)
# Confusion Matrix
conf_matrix = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:\n", conf_matrix)
# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

# Precision, Recall, F1-score
precision, recall, f1_score, _ = precision_recall_fscore_support(true_classes,
    ↪ predicted_classes, average='weighted')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1_score:.4f}')

```

275/275 [=====] - 1s 3ms/step - loss: 0.5734 -

accuracy: 0.8130

Test Accuracy: 81.30%

Test Loss: 0.5734

Classification Report:


	precision	recall	f1-score	support
0	0.83	0.86	0.84	3129
1	0.90	0.86	0.88	1466
2	0.77	0.81	0.79	1784
3	0.86	0.63	0.73	357
4	0.76	0.74	0.75	2064
accuracy			0.81	8800
macro avg	0.82	0.78	0.80	8800
weighted avg	0.81	0.81	0.81	8800

Confusion Matrix:

```

[[2690  19 137   9 274]
 [ 84 1262  60  12  48]
 [ 134  67 1446   2 135]
 [ 56  32  16 225  28]
 [ 289  21  208  15 1531]]

```



Precision: 0.8139

Recall: 0.8130

F1-score: 0.8125

```
[ ]: # Extract TP, TN, FP, FN from confusion matrix
```

```
TP = conf_matrix[1, 1] # True Positives
```

```
TN = conf_matrix[0, 0] # True Negatives
```

```
FP = conf_matrix[0, 1] # False Positives
```

```
FN = conf_matrix[1, 0] # False Negatives
```

```
print("True Positives:", TP)
```

```
print("True Negatives:", TN)
```

```
print("False Positives:", FP)
```

```
print("False Negatives:", FN)
```

True Positives: 1262

True Negatives: 2690

False Positives: 19

False Negatives: 84

```
[ ]: # Save the entire model to a HDF5 file
model.save('/content/drive/MyDrive/Bully/CNN_Model')
```

Ensemble Model

```
[ ]: models = {'GRU': GRUmodel(tokenizer.vocab_size),
               'LSTM': LSTMmodel(tokenizer.vocab_size),
               'CNN': CNNmodel(tokenizer.vocab_size),
               'HybridModel' : HybridModel(tokenizer.vocab_size)}
```

```
[ ]: histories = {} # To store training histories of each model

for name, model in models.items():
    print(f"Training {name} model...")
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    history = model.fit(train_encodings['input_ids'], train_labels_onehot,
epochs=5, batch_size=32)
    histories[name] = history
```

Training GRU model...

Epoch 1/5

1100/1100 [=====] - 92s 75ms/step - loss: 1.0465 - accuracy: 0.5912

Epoch 2/5

1100/1100 [=====] - 36s 33ms/step - loss: 0.6605 - accuracy: 0.7933

Epoch 3/5

1100/1100 [=====] - 37s 34ms/step - loss: 0.5180 - accuracy: 0.8407

Epoch 4/5

1100/1100 [=====] - 34s 31ms/step - loss: 0.4455 - accuracy: 0.8665

Epoch 5/5

1100/1100 [=====] - 45s 41ms/step - loss: 0.3802 - accuracy: 0.8856

Training LSTM model...

Epoch 1/5

1100/1100 [=====] - 87s 70ms/step - loss: 0.9218 - accuracy: 0.6537

Epoch 2/5

1100/1100 [=====] - 33s 30ms/step - loss: 0.5596 - accuracy: 0.8273

Epoch 3/5

1100/1100 [=====] - 27s 24ms/step - loss: 0.4517 - accuracy: 0.8610

Epoch 4/5

1100/1100 [=====] - 24s 22ms/step - loss: 0.3739 -

```

accuracy: 0.8837
Epoch 5/5
1100/1100 [=====] - 23s 21ms/step - loss: 0.3132 -
accuracy: 0.9013
Training CNN model...
Epoch 1/5
1100/1100 [=====] - 40s 35ms/step - loss: 0.7682 -
accuracy: 0.7190
Epoch 2/5
1100/1100 [=====] - 13s 12ms/step - loss: 0.4756 -
accuracy: 0.8467
Epoch 3/5
1100/1100 [=====] - 12s 11ms/step - loss: 0.3416 -
accuracy: 0.8912
Epoch 4/5
1100/1100 [=====] - 9s 9ms/step - loss: 0.2493 -
accuracy: 0.9231
Epoch 5/5
1100/1100 [=====] - 11s 10ms/step - loss: 0.1784 -
accuracy: 0.9451
Training HybridModel model...
Epoch 1/5
1100/1100 [=====] - 88s 68ms/step - loss: 1.2017 -
accuracy: 0.4947
Epoch 2/5
1100/1100 [=====] - 51s 47ms/step - loss: 0.7415 -
accuracy: 0.7593
Epoch 3/5
1100/1100 [=====] - 48s 43ms/step - loss: 0.5731 -
accuracy: 0.8186
Epoch 4/5
1100/1100 [=====] - 48s 43ms/step - loss: 0.4760 -
accuracy: 0.8562
Epoch 5/5
1100/1100 [=====] - 47s 42ms/step - loss: 0.4109 -
accuracy: 0.8775

```

```
[ ]: # Plotting training accuracy and loss for each model
```

```

plt.figure(figsize=(12, 6))

for name, history in histories.items():
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label=f'{name} Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.title('Training Accuracy')
    plt.legend()


```

```

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label=f'{name} Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss')
plt.legend()

plt.tight_layout()
plt.show()

```



```

[ ]: train_predictions = {}
test_predictions = {}

for name, model in models.items():
    train_predictions[name] = model.predict(train_encodings['input_ids'])
    test_predictions[name] = model.predict(test_encodings['input_ids'])

```

```

1100/1100 [=====] - 11s 10ms/step
275/275 [=====] - 3s 10ms/step
1100/1100 [=====] - 9s 8ms/step
275/275 [=====] - 2s 7ms/step
1100/1100 [=====] - 2s 2ms/step
275/275 [=====] - 1s 2ms/step
1100/1100 [=====] - 18s 16ms/step
275/275 [=====] - 4s 15ms/step

```

```
[ ]: stacked_train_predictions = np.column_stack([train_predictions[name] for name in models])
stacked_test_predictions = np.column_stack([test_predictions[name] for name in models])
```

```
[ ]: stacked_train, stacked_val, train_labels_train, train_labels_val = train_test_split(
    stacked_train_predictions,
    train_labels_onehot,
    test_size=0.2,
    random_state=42
)
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: # Initialize the Random Forest model
rf_meta_learner = RandomForestClassifier(n_estimators=100, random_state=42)
rf_meta_learner.fit(stacked_train, train_labels_train)
rf_meta_predictions = rf_meta_learner.predict(stacked_val)
```

```
[ ]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(train_labels_val, rf_meta_predictions)
print(f"Accuracy of Random Forest meta-learner: {accuracy * 100:.2f}%")
```

Accuracy of Random Forest meta-learner: 96.38%

```
[ ]: from sklearn.metrics import precision_recall_fscore_support

precision, recall, f1_score, _ = precision_recall_fscore_support(
    train_labels_val, rf_meta_predictions,
    average='weighted')

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1_score:.4f}")
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	2483
1	0.97	0.99	0.98	1209
2	0.97	0.98	0.97	1452
3	0.92	0.91	0.91	269
4	0.96	0.94	0.95	1627
accuracy			0.97	7040

macro avg	0.96	0.96	0.96	7040
weighted avg	0.97	0.97	0.97	7040

```
[ ]: from sklearn.metrics import precision_recall_fscore_support
```

```
[ ]: # Extract precision, recall, and f1-score
precision, recall, f1_score, _ = precision_recall_fscore_support(true_labels,
    ↪ log_reg_predictions, average='weighted')

# Print the results
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1_score:.4f}")
```

```
Precision: 0.9704
Recall: 0.9705
F1 Score: 0.9704
```

```
[ ]: from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_curve, auc
```

```
[ ]: from sklearn.preprocessing import label_binarize

if len(train_labels_val.shape) > 1 and train_labels_val.shape[1] > 1:
    train_labels_val = np.argmax(train_labels_val, axis=1)

# Binarize the labels for multiclass ROC calculation
label_binarizer = label_binarize(train_labels_val, classes=np.
    ↪ unique(train_labels_val))


# Calculate ROC curve for the Random Forest meta-learner
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(label_binarizer.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(label_binarizer[:, i], rf_meta_predictions[:,
    ↪ i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC curve for each class
plt.figure(figsize=(8, 6))

for i in range(label_binarizer.shape[1]):
    plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')
```

```
plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Each Class')
plt.legend(loc="lower right")
plt.show()
```



```
[ ]: # Binarize the labels for multiclass ROC calculation
label_binarizer = label_binarize(train_labels_val, classes=np.
    ↪unique(train_labels_val))

# Ensure the number of samples match
if label_binarizer.shape[0] != log_reg_predictions.shape[0]:
    raise ValueError("Inconsistent number of samples between label_binarizer_
    ↪and log_reg_predictions.")

# Calculate ROC curve for the Logistic Regression meta-learner
fpr, tpr, _ = roc_curve(label_binarizer.ravel(), log_reg_predictions.ravel())
roc_auc = auc(fpr, tpr)
```



```

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-111-61b1966bb40c> in <cell line: 9>()
      7
      8 # Calculate ROC curve for the Logistic Regression meta-learner
----> 9 fpr, tpr, _ = roc_curve(label_binarizer.ravel(), log_reg_predictions.
    ↪ ravel())
     10 roc_auc = auc(fpr, tpr)
     11

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_ranking.py in
    ↪ roc_curve(y_true, y_score, pos_label, sample_weight, drop_intermediate)
     990     array([1.8 , 0.8 , 0.4 , 0.35, 0.1 ])
     991     """
--> 992     fps, tps, thresholds = _binary_clf_curve(
     993         y_true, y_score, pos_label=pos_label, sample_weight=sample_weight
     994     )

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_ranking.py in
    ↪ _binary_clf_curve(y_true, y_score, pos_label, sample_weight)
     749         raise ValueError("{0} format is not supported".format(y_type))
     750
--> 751     check_consistent_length(y_true, y_score, sample_weight)
     752     y_true = column_or_1d(y_true)
     753     y_score = column_or_1d(y_score)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in
    ↪ check_consistent_length(*arrays)
     395     uniques = np.unique(lengths)
     396     if len(uniques) > 1:
--> 397         raise ValueError(
     398             "Found input variables with inconsistent numbers of samples
    ↪ %r"
     399             % [int(l) for l in lengths]

```

```
ValueError: Found input variables with inconsistent numbers of samples: [35200,
↪7040]
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: !pip install joblib
```

Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (1.3.2)

```
[ ]: from joblib import dump
```

```
# Assuming rf_meta_learner is your trained Random Forest meta-learner
model_filename = '/content/drive/MyDrive/Bully/rf_meta_learner.joblib'

# Save the trained model to a file
dump(rf_meta_learner, model_filename)
```

```
[ ]: #from sklearn.preprocessing import label_binarize
#
#if len(train_labels_val.shape) > 1 and train_labels_val.shape[1] > 1:
#    train_labels_val = np.argmax(train_labels_val, axis=1)
#
## Binarize the labels for multiclass ROC calculation
#label_binarizer = label_binarize(train_labels_val, classes=np.
↪unique(train_labels_val))
#
## Calculate ROC curve for the Random Forest meta-learner
#fpr = dict()
#tpr = dict()
#roc_auc = dict()
#
#for i in range(label_binarizer.shape[1]):
#    fpr[i], tpr[i], _ = roc_curve(label_binarizer[:, i], rf_meta_predictions[:
↪, i])
#    roc_auc[i] = auc(fpr[i], tpr[i])
#
```

```

## Plot ROC curve for each class
plt.figure(figsize=(8, 6))
#
#for i in range(label_binarizer.shape[1]):
#     plt.plot(fpr[i], tpr[i], label=f'Class {i} (AUC = {roc_auc[i]:.2f})')
#
plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Each Class')
plt.legend(loc="lower right")
plt.show()

```

```

[ ]: from sklearn.preprocessing import LabelEncoder

if len(train_labels_val.shape) > 1 and train_labels_val.shape[1] > 1:
    train_labels_val = np.argmax(train_labels_val, axis=1)

# Use LabelEncoder to ensure train_labels_val is represented as integers
label_encoder = LabelEncoder()
train_labels_val_encoded = label_encoder.fit_transform(train_labels_val)

# Convert rf_meta_predictions to integer classes
rf_meta_predictions_int = np.argmax(rf_meta_predictions, axis=1)

# Ensure shapes align
print("Shapes - True Labels:", train_labels_val_encoded.shape, "Predictions:",
      rf_meta_predictions_int.shape)

# Then try calculating confusion matrix and classification report
conf_matrix = confusion_matrix(train_labels_val_encoded,
                               rf_meta_predictions_int)
class_report = classification_report(train_labels_val_encoded,
                                     rf_meta_predictions_int)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)

```

```

[ ]: # Plotting the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='d', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```

Lime

```
[ ]: !pip install shap lime
```

Collecting shap

Downloading shap-0.44.0-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64
.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (533 kB)
533.5/533.5

kB 8.3 MB/s eta 0:00:00

Collecting lime

Downloading lime-0.2.0.1.tar.gz (275 kB)
275.7/275.7

kB 35.5 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages
(from shap) (1.23.5)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages
(from shap) (1.11.4)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-
packages (from shap) (1.2.2)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages
(from shap) (1.5.3)

Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-
packages (from shap) (4.66.1)

Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-
packages (from shap) (23.2)

Collecting slicer==0.0.7 (from shap)

Downloading slicer-0.0.7-py3-none-any.whl (14 kB)

Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages
(from shap) (0.58.1)

Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-
packages (from shap) (2.2.1)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-
packages (from lime) (3.7.1)

Requirement already satisfied: scikit-image>=0.12 in
/usr/local/lib/python3.10/dist-packages (from lime) (0.19.3)

Requirement already satisfied: networkx>=2.2 in /usr/local/lib/python3.10/dist-
packages (from scikit-image>=0.12->lime) (3.2.1)

Requirement already satisfied: pillow!=7.1.0,!7.1.1,!8.3.0,>=6.1.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (9.4.0)

Requirement already satisfied: imageio>=2.4.1 in /usr/local/lib/python3.10/dist-
packages (from scikit-image>=0.12->lime) (2.31.6)

Requirement already satisfied: tifffile>=2019.7.26 in
/usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime)
(2023.12.9)

Requirement already satisfied: PyWavelets>=1.1.1 in
/usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (1.5.0)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-

```

packages (from scikit-learn->shap) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.2.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->lime) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-
packages (from matplotlib->lime) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->lime) (4.47.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->lime) (1.4.5)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->lime) (3.1.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->lime) (2.8.2)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in
/usr/local/lib/python3.10/dist-packages (from numba->shap) (0.41.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas->shap) (2023.3.post1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.7->matplotlib->lime) (1.16.0)
Building wheels for collected packages: lime
  Building wheel for lime (setup.py) ... done
  Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.whl size=283835
sha256=57526409ff2e0b8024478bfe93d6d1aa7a8ad908d7ded4b4cbe985f1f18ac94f
  Stored in directory: /root/.cache/pip/wheels/fd/a2/af/9ac0a1a85a27f314a06b39e1
f492bee1547d52549a4606ed89
Successfully built lime
Installing collected packages: slicer, shap, lime
Successfully installed lime-0.2.0.1 shap-0.44.0 slicer-0.0.7

```

```
[ ]: import shap
      from lime import lime_tabular
```

```
[ ]: !free -h
```

	total	used	free	shared	buff/cache	available
Mem:	12Gi	3.9Gi	2.9Gi	21Mi	5.9Gi	8.5Gi
Swap:	0B	0B	0B			

```
[ ]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[ ]: X = df['comment']
      y = df['label']

      # Split the data into training and validation sets
```

```
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,  
↪random_state=42)
```

```
# Vectorize the text data using TF-IDF  
vectorizer = TfidfVectorizer()  
X_train_vec = vectorizer.fit_transform(X_train)  
X_val_vec = vectorizer.transform(X_val)
```

```
[ ]: # Get the unique topic names from the 'label' column  
topic_names = df['label'].unique()  
print("Unique topic names:", topic_names)
```

Unique topic names: [2 0 4 1 3]

```
[ ]: !pip install textblob
```

Requirement already satisfied: textblob in /usr/local/lib/python3.10/dist-packages (0.17.1)
Requirement already satisfied: nltk>=3.1 in /usr/local/lib/python3.10/dist-packages (from textblob) (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk>=3.1->textblob) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk>=3.1->textblob) (1.3.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk>=3.1->textblob) (2023.6.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk>=3.1->textblob) (4.66.1)

```
[ ]: from textblob import TextBlob  
  
def analyze_sentiment(text):  
    # Create a TextBlob object  
    blob = TextBlob(text)  
  
    # Perform sentiment analysis  
    sentiment_score = blob.sentiment.polarity  
  
    return sentiment_score
```

```
[ ]: !pip install gensim
```

Requirement already satisfied: gensim in /usr/local/lib/python3.10/dist-packages (4.3.2)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.10/dist-packages (from gensim) (1.23.5)
Requirement already satisfied: scipy>=1.7.0 in /usr/local/lib/python3.10/dist-packages (from gensim) (1.11.4)

Requirement already satisfied: smart-open>=1.8.1 in
/usr/local/lib/python3.10/dist-packages (from gensim) (6.4.0)

```
[ ]: from gensim import corpora, models
import matplotlib.pyplot as plt

# Assuming 'corpus' is a preprocessed list of text documents

# Create a dictionary representation of the documents
dictionary = corpora.Dictionary(corpus)

# Convert the corpus to a document-term matrix
doc_term_matrix = [dictionary.doc2bow(doc) for doc in corpus]

# Train the LDA model
lda_model = models.LdaModel(doc_term_matrix, num_topics=5, id2word=dictionary,
    ↪passes=15)

# Assign labels/categories based on the identified topics
topics = lda_model.show_topics(formatted=False)
topic_labels = {0: 'not bully', 1: 'troll', 2: 'sexual', 3: 'religious', 4:
    ↪'threat'} # Assign labels manually

# Plotting the probabilities of categories
for topic_id, topic in topics:
    topic_probabilities = [prob for _, prob in topic]
    plt.plot(topic_probabilities, label=topic_labels[topic_id])

plt.xlabel('Word ID')
plt.ylabel('Probability')
plt.legend()
plt.show()
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-146-6f4b67661efe> in <cell line: 7>()
      5
      6 # Create a dictionary representation of the documents
----> 7 dictionary = corpora.Dictionary(corpus)
      8
      9 # Convert the corpus to a document-term matrix

/usr/local/lib/python3.10/dist-packages/gensim/corpora/dictionary.py in
    ↪__init__(self, documents, prune_at)
      76
      77         if documents is not None:
----> 78             self.add_documents(documents, prune_at=prune_at)
```

```

79         self.add_lifecycle_event(
80             "created",

/usr/local/lib/python3.10/dist-packages/gensim/corpora/dictionary.py in
↳ add_documents(self, documents, prune_at)
202
203         # update Dictionary with the document
--> 204         self.doc2bow(document, allow_update=True) # ignore the
↳ result, here we only care about updating token ids
205
206         logger.info("built %s from %i documents (total %i corpus
↳ positions)", self, self.num_docs, self.num_pos)

/usr/local/lib/python3.10/dist-packages/gensim/corpora/dictionary.py in
↳ doc2bow(self, document, allow_update, return_missing)
239         """
240         if isinstance(document, str):
--> 241             raise TypeError("doc2bow expects an array of unicode tokens,
↳ on input, not a single string")
242
243         # Construct (word, frequency) mapping.

TypeError: doc2bow expects an array of unicode tokens on input, not a single
↳ string

```

```

[ ]: # Sample input text
input_text = "

↳

# Tokenize and preprocess the input text
processed_input = dictionary.doc2bow(simple_preprocess(input_text))

# Get the topic distribution for the input text
topic_distribution = lda_model[processed_input]

# Display the extracted topics and their probabilities for the input text
print(topic_distribution)

```

```

[(0, 0.018206181), (1, 0.018206127), (2, 0.018206157), (3, 0.01820615), (4,
0.9271754)]

```

```

[ ]:

```

```

[ ]: from transformers import BertTokenizer

# Load the Bangla BERT tokenizer
tokenizer = BertTokenizer.from_pretrained("sagorsarker/bangla-bert-base")

```



```

# Input text
input_text = "

# Tokenize the input text using the BERT tokenizer
tokenized_input = tokenizer.encode_plus(
    input_text,
    max_length=100,
    truncation=True,
    padding="max_length",
    return_tensors="np"
)

# Extract features
word_count = len(tokenized_input["input_ids"][0])
sentence_length = len(input_text.split('.'))
average_word_length = sum(len(word) for word in input_text.split()) / word_count

# Sentiment analysis (this requires a trained sentiment analysis model)
sentiment_score = analyze_sentiment(input_text) # Placeholder function

# Assuming you need BERT token indices as features
bert_token_indices = tokenized_input["input_ids"][0]

# Display extracted features
print(f"Word count: {word_count}")
print(f"Sentence length: {sentence_length}")
print(f"Average word length: {average_word_length}")
print(f"Sentiment score: {sentiment_score}")
print(f"BERT token indices: {bert_token_indices}")

```

Word count: 100

Sentence length: 1

Average word length: 1.28

Sentiment score: 0.0

BERT token indices: [101 10706 7448 4931 85209 9294 2844 25151 35705 5931
5740 9294

9294	51264	6741	72612	5843	7932	2093	100	2285	45931	2097	1014
16137	21327	2040	6741	33983	59464	8833	2237	7373	1011	16137	21327
2040	4374	2076	5500	6399	3364	5740	9294	9294	2069	8705	9294
2446	21327	5740	9294	9294	2058	2069	3447	5843	1014	102	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0]								

```
[ ]: encoded_text = tokenizer.encode(input_text, max_length=100, truncation=True,
    ↪padding='max_length', return_tensors='pt')
bert_tokens = encoded_text[0].tolist() # Converting tensor to list
print("BERT token indices:", bert_tokens)
```

```
BERT token indices: [101, 10706, 7448, 4931, 85209, 9294, 2844, 25151, 35705,
5931, 5740, 9294, 9294, 51264, 6741, 72612, 5843, 7932, 2093, 100, 2285, 45931,
2097, 1014, 16137, 21327, 2040, 6741, 33983, 59464, 8833, 2237, 7373, 1011,
16137, 21327, 2040, 4374, 2076, 5500, 6399, 3364, 5740, 9294, 9294, 2069, 8705,
9294, 2446, 21327, 5740, 9294, 9294, 2058, 2069, 3447, 5843, 1014, 102, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
[ ]: # Define the feature names
feature_names = []

# BERT token indices
for i in range(100): # Assuming max_length=100 for BERT tokenizer
    feature_names.append(f"token_{i+1}")

# Text statistics
feature_names.extend([
    'word_count',
    'sentence_length',
    'average_word_length',
])
```

```
[ ]: # Sentiment analysis score
feature_names.append('sentiment_score')

# Model-specific features
for model_name in ['BiGRU', 'BiLSTM', 'CNN']:
    for i in range(50): # Assuming 50 intermediate outputs for each model
        feature_names.append(f"{model_name}_output_{i+1}")
```

```
[ ]: # Meta-model features (logistic regression)
for i in range(10): # Assuming 10 features in the logistic regression
    ↪meta-model
    feature_names.append(f"log_reg_feature_{i+1}")

# Generate random topic probabilities for demonstration
num_samples = len(stacked_val)
num_topics = 5
topic_probabilities = np.random.rand(num_samples, num_topics)


# Example feature names for topic probabilities
for i in range(num_topics):
```

```
feature_names.append(f"topic_{i+1}_probability")
```

```
[ ]: # Use SHAP for explanation
explainer = shap.Explainer(log_reg, stacked_train)
shap_values = explainer.shap_values(stacked_val)


# Assuming 'label_names' is a list of class names for classification
label_names = [
    'not bully',
    'troll',
    'sexual',
    'religious',
    'threat'
]

# Visualize SHAP summary plot
shap.summary_plot(shap_values, stacked_val, feature_names=feature_names,
↳class_names=label_names)
```



```
[ ]: # Using SHAP for explanation
explainer = shap.Explainer(log_reg, stacked_train) # log_reg is your logistic
↪ regression model
shap_values = explainer.shap_values(stacked_val)

# Visualize SHAP summary plot
shap.summary_plot(shap_values, stacked_val, feature_names=feature_names,
↪ class_names=label_names)
```



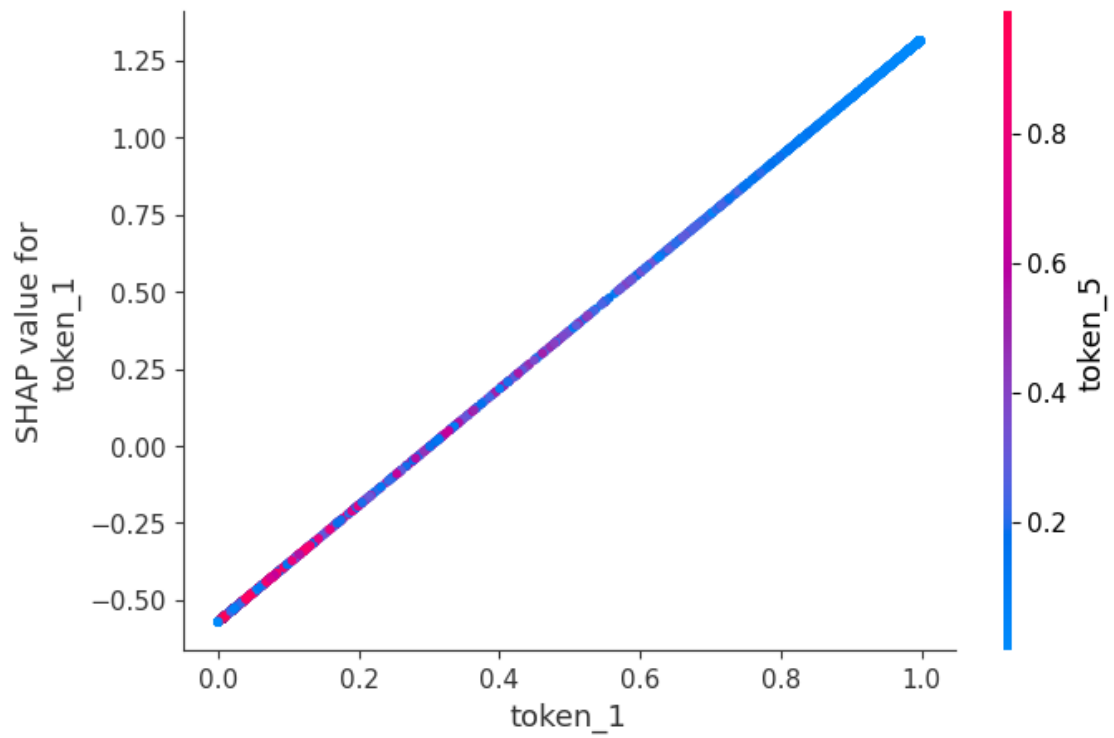
```
[ ]: print(stacked_val.shape)
```

```
(7040, 15)
```


```
[ ]: print(len(feature_names))
```

```
269
```

```
[ ]: # Replace '0' with the desired feature index
shap.dependence_plot(0, shap_values[0], stacked_val,
↪ feature_names=feature_names)
```




```
[ ]: # Replace '0' with the desired feature index
shap.dependence_plot(0, shap_values[0], stacked_val,
    ↪ feature_names=feature_names, interaction_index='auto')
```



```
[ ]: # Replace 'label_names' with your actual label names
label_names = [
    'not bully',
    'troll',
    'sexual',
    'religious',
    'threat'
]

# For summary plot
shap.summary_plot(shap_values, stacked_val, feature_names=feature_names,
                  ↪class_names=label_names)
```



```
[ ]: print("Size of SHAP values:", len(shap_values))
      print("Size of feature names:", len(feature_names))
```

Size of SHAP values: 5
Size of feature names: 269

```
[ ]: class_index = 0 # Replace this with the desired class index (0 to 14)
      shap.force_plot(explainer.expected_value[class_index],
                      ↪shap_values[class_index][0], stacked_val[0], feature_names=feature_names)
```

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-171-100139b251ae> in <cell line: 2>()
    1 class_index = 0 # Replace this with the desired class index (0 to 14)
```



```

----> 2 shap.force_plot(explainer.expected_value[class_index],
↳shap_values[class_index][0], stacked_val[0], feature_names=feature_names)
    3

/usr/local/lib/python3.10/dist-packages/shap/plots/_force.py in
↳force(base_value, shap_values, features, feature_names, out_names, link,
↳plot_cmap, matplotlib, show, figsize, ordering_keys,
↳ordering_keys_time_format, text_rotation, contribution_threshold)
    194     )
    195
--> 196     return visualize(e,

    197         plot_cmap,
    198         matplotlib,

/usr/local/lib/python3.10/dist-packages/shap/plots/_force.py in visualize(e,
↳plot_cmap, matplotlib, figsize, show, ordering_keys,
↳ordering_keys_time_format, text_rotation, min_perc)
    412     )
    413     else:
--> 414         return AdditiveForceVisualizer(e, plot_cmap=plot_cmap)
    415     elif isinstance(e, Explanation):
    416         if matplotlib:

/usr/local/lib/python3.10/dist-packages/shap/plots/_force.py in __init__(self,
↳e, plot_cmap)
    490         # build the json data
    491         features = {}
--> 492         for i in filter(lambda j: e.effects[j] != 0, range(len(e.data.
↳group_names))):
    493             features[i] = {
    494                 "effect": ensure_not_numpy(e.effects[i]),

/usr/local/lib/python3.10/dist-packages/shap/plots/_force.py in <lambda>(j)
    490         # build the json data
    491         features = {}
--> 492         for i in filter(lambda j: e.effects[j] != 0, range(len(e.data.
↳group_names))):
    493             features[i] = {
    494                 "effect": ensure_not_numpy(e.effects[i]),

IndexError: index 15 is out of bounds for axis 0 with size 15

```

```
[ ]:
```

```
[ ]: num_samples = 1000
num_topics = 5
```

```
topic_probabilities = np.random.rand(num_samples, num_topics)
```

```
# Example feature names for topic probabilities
```

```
feature_names = []
```

```
for i in range(num_topics):
```

```
    feature_names.append(f"topic_{i+1}_probability")
```

```
[ ]: # Create a SHAP explainer object for the logistic regression model
```

```
explainer = shap.Explainer(log_reg, stacked_train)
```

```
# Calculate SHAP values for the validation data
```

```
shap_values = explainer.shap_values(stacked_val)
```

```
# Visualize SHAP summary plot
```

```
shap.summary_plot(shap_values, stacked_val, feature_names=feature_names)
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: # Create a SHAP explainer object for the logistic regression model
```


```
explainer = shap.Explainer(log_reg, stacked_train)
```

```
# Calculate SHAP values for the validation data
```

```
shap_values = explainer.shap_values(stacked_val)
```

```
# Visualize SHAP summary plot
```

```
shap.summary_plot(shap_values, stacked_val)
```



```
[ ]: num_features = stacked_val.shape[1]
print("Number of features in stacked_val:", num_features)
print("Number of elements in label_names:", len(label_names))
```

Number of features in stacked_val: 15
Number of elements in label_names: 5

```
[ ]: def predict_log_reg_proba(texts):
    # Assuming you have a tokenizer for text data
    encodings = tokenizer(texts, truncation=True, padding='max_length',
    ↪max_length=100, return_tensors='tf')['input_ids']
    inputs = {'sequences': encodings}

    # Use your 'log_reg' Logistic Regression model to predict probabilities
```

```

    predictions = log_reg.predict_proba(inputs) # Replace 'inputs' with the
↳appropriately preprocessed data

    return predictions

```

```

[ ]: feature_names = [
    'not bully',
    'troll',
    'sexual',
    'religious',
    'threat'
]

```

```

[ ]: from lime.lime_tabular import LimeTabularExplainer

# Assuming stacked_val represents your validation data for the Logistic
↳Regression model
# Replace 'feature_names' with the actual names of the features used in the
↳Logistic Regression model
# 'train_labels_val' should be the labels corresponding to 'stacked_val'
explainer = LimeTabularExplainer(np.array(stacked_val),
                                mode='classification',
                                training_labels=train_labels_val,
                                feature_names=feature_names, # Replace with
↳actual feature names

                                discretize_continuous=True)

```

```

-----
IndexError                                Traceback (most recent call last)
<ipython-input-123-9477b49d77f7> in <cell line: 6>()
      4 # Replace 'feature_names' with the actual names of the features used in,
↳the Logistic Regression model
      5 # 'train_labels_val' should be the labels corresponding to 'stacked_val'
----> 6 explainer = LimeTabularExplainer(np.array(stacked_val),
      7
      8                                mode='classification',
                                training_labels=train_labels_val,

/usr/local/lib/python3.10/dist-packages/lime/lime_tabular.py in __init__(self,
↳training_data, mode, training_labels, feature_names, categorical_features,
↳categorical_names, kernel_width, kernel, verbose, class_names,
↳feature_selection, discretize_continuous, discretizer, sample_around_instance
↳random_state, training_data_stats)
    213
    214         if discretizer == 'quartile':
--> 215             self.discretizer = QuartileDiscretizer(
    216
    217                 training_data, self.categorical_features,
                                self.feature_names, labels=training_labels,

```

```

/usr/local/lib/python3.10/dist-packages/lime/discretize.py in __init__(self,
↳data, categorical_features, feature_names, labels, random_state)
    176     def __init__(self, data, categorical_features, feature_names,
↳labels=None, random_state=None):
    177
--> 178         BaseDiscretizer.__init__(self, data, categorical_features,
    179                                   feature_names, labels=labels,
    180                                   random_state=random_state)

/usr/local/lib/python3.10/dist-packages/lime/discretize.py in __init__(self,
↳data, categorical_features, feature_names, labels, random_state, data_stats)
    62         n_bins = qts.shape[0] # Actually number of borders (=
↳#bins-1)
    63         boundaries = np.min(data[:, feature]), np.max(data[:,
↳feature])
---> 64         name = feature_names[feature]
    65
    66         self.names[feature] = ['%s <= %.2f' % (name, qts[0])]

IndexError: list index out of range

```

```
[ ]:
```

```
[ ]:
```

```
[ ]: # Initialize KernelExplainer
kernel_explainer = shap.KernelExplainer(rf_meta_learner.predict, stacked_train)
```

WARNING:shap:Using 28158 background data samples could cause slower run times. Consider using shap.sample(data, K) or shap.kmeans(data, K) to summarize the background as K samples.

```
[ ]: import shap

# Initialize KernelExplainer
kernel_explainer = shap.KernelExplainer(rf_meta_learner.predict, stacked_train)

# Sample a subset for SHAP visualization (adjust the subset size based on your
↳resources)
subset_size_shap = 20
subset_indices_shap = np.random.choice(len(test_texts), size=subset_size_shap,
↳replace=False)
subset_test_texts = [test_texts[i] for i in subset_indices_shap]
subset_test_input_ids = tokenizer(subset_test_texts, truncation=True,
↳padding=True, max_length=max_length, return_tensors="tf")

```

```
subset_stacked_val_embeddings = models['GRU'].
    ↪predict(subset_test_input_ids['input_ids'].numpy())
```

WARNING:shap:Using 28158 background data samples could cause slower run times. Consider using shap.sample(data, K) or shap.kmeans(data, K) to summarize the background as K samples.

1/1 [=====] - 2s 2s/step

```
[ ]: import shap
import numpy as np

# Sample a subset for SHAP visualization (adjust the subset size based on your
    ↪resources)
subset_size_shap = 20
subset_indices_shap = np.random.choice(len(test_texts), size=subset_size_shap,
    ↪replace=False)
subset_test_texts = [test_texts[i] for i in subset_indices_shap]
subset_test_input_ids = tokenizer(subset_test_texts, truncation=True,
    ↪padding=True, max_length=max_length, return_tensors="tf")
subset_stacked_val_embeddings = models['GRU'].
    ↪predict(subset_test_input_ids['input_ids'].numpy())

# Initialize KernelExplainer with the RandomForest model
kernel_explainer = shap.KernelExplainer(rf_meta_learner.predict, shap.
    ↪sample(stacked_train, 100))

# Calculate SHAP values using KernelExplainer
shap_values_kernel = kernel_explainer.shap_values(subset_stacked_val_embeddings)

# Calculate SHAP values using KernelExplainer
shap_values_kernel = kernel_explainer.shap_values(subset_stacked_val_embeddings)

# Summary plot for the first instance
shap.summary_plot(shap_values_kernel, subset_stacked_val_embeddings,
    ↪feature_names=models.keys())
```

1/1 [=====] - 0s 37ms/step

0%| | 0/20 [00:00<?, ?it/s]

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-103-98bbc0632a55> in <cell line: 15>()
    13
    14 # Calculate SHAP values using KernelExplainer
--> 15 shap_values_kernel = kernel_explainer.
    ↪shap_values(subset_stacked_val_embeddings)
```

```

16
17 # Calculate SHAP values using KernelExplainer

/usr/local/lib/python3.10/dist-packages/shap/explainers/_kernel.py in
↳shap_values(self, X, **kwargs)
    242         if self.keep_index:
    243             data = convert_to_instance_with_index(data,
↳column_name, index_value[i:i + 1], index_name)
--> 244             explanations.append(self.explain(data, **kwargs))
    245             if kwargs.get("gc_collect", False):
    246                 gc.collect()

/usr/local/lib/python3.10/dist-packages/shap/explainers/_kernel.py in
↳explain(self, incoming_instance, **kwargs)
    269         # convert incoming input to a standardized iml object
    270         instance = convert_to_instance(incoming_instance)
--> 271         match_instance_to_data(instance, self.data)
    272
    273         # find the feature groups we will test. If a feature does not
↳change from its

/usr/local/lib/python3.10/dist-packages/shap/utils/_legacy.py in
↳match_instance_to_data(instance, data)
    88         if isinstance(data, DenseData):
    89             if instance.group_display_values is None:
--> 90                 instance.group_display_values = [instance.x[0, group[0]] if
↳len(group) == 1 else "" for group in data.groups]
    91         assert len(instance.group_display_values) == len(data.groups)
    92         instance.groups = data.groups

/usr/local/lib/python3.10/dist-packages/shap/utils/_legacy.py in <listcomp>(.0)
    88         if isinstance(data, DenseData):
    89             if instance.group_display_values is None:
--> 90                 instance.group_display_values = [instance.x[0, group[0]] if
↳len(group) == 1 else "" for group in data.groups]
    91         assert len(instance.group_display_values) == len(data.groups)
    92         instance.groups = data.groups

IndexError: index 5 is out of bounds for axis 1 with size 5

```

```

[ ]: print("Shape of subset_stacked_val_embeddings:", subset_stacked_val_embeddings.
↳shape)
print("Shape of background dataset (stacked_train):", stacked_train.shape)
shap_values_kernel = kernel_explainer.shap_values(subset_stacked_val_embeddings)

```

```

Shape of subset_stacked_val_embeddings: (20, 5)
Shape of background dataset (stacked_train): (28158, 15)

```

```

-----
IndexError                                Traceback (most recent call last)
<ipython-input-91-71190363c975> in <cell line: 3>()
      1 print("Shape of subset_stacked_val_embeddings:",
↳ subset_stacked_val_embeddings.shape)
      2 print("Shape of background dataset (stacked_train):", stacked_train.
↳ shape)
----> 3 shap_values_kernel = kernel_explainer.
↳ shap_values(subset_stacked_val_embeddings)

/usr/local/lib/python3.10/dist-packages/shap/explainers/_kernel.py in
↳ shap_values(self, X, **kwargs)
    242         if self.keep_index:
    243             data = convert_to_instance_with_index(data,
↳ column_name, index_value[i:i + 1], index_name)
--> 244             explanations.append(self.explain(data, **kwargs))
    245             if kwargs.get("gc_collect", False):
    246                 gc.collect()

/usr/local/lib/python3.10/dist-packages/shap/explainers/_kernel.py in
↳ explain(self, incoming_instance, **kwargs)
    269         # convert incoming input to a standardized iml object
    270         instance = convert_to_instance(incoming_instance)
--> 271         match_instance_to_data(instance, self.data)
    272
    273         # find the feature groups we will test. If a feature does not
↳ change from its

/usr/local/lib/python3.10/dist-packages/shap/utils/_legacy.py in
↳ match_instance_to_data(instance, data)
    88         if isinstance(data, DenseData):
    89             if instance.group_display_values is None:
--> 90                 instance.group_display_values = [instance.x[0, group[0]] if
↳ len(group) == 1 else "" for group in data.groups]
    91         assert len(instance.group_display_values) == len(data.groups)
    92         instance.groups = data.groups

/usr/local/lib/python3.10/dist-packages/shap/utils/_legacy.py in <listcomp>(.0)
    88         if isinstance(data, DenseData):
    89             if instance.group_display_values is None:
--> 90                 instance.group_display_values = [instance.x[0, group[0]] if
↳ len(group) == 1 else "" for group in data.groups]
    91         assert len(instance.group_display_values) == len(data.groups)
    92         instance.groups = data.groups

```



```
IndexError: index 5 is out of bounds for axis 1 with size 5
```

```
[ ]: # Print intermediate shapes and values for debugging
print("Shape of BERT embeddings (before):", subset_stacked_val_embeddings.shape)

# Check the structure of the obtained embeddings
print("First instance of BERT embeddings:", subset_stacked_val_embeddings[0])
```

Shape of BERT embeddings (before): (10, 5)

First instance of BERT embeddings: [0.00691598 0.01117767 0.88472986 0.0036564
0.09352009]

```
[ ]: # Check the shape of the model training data
print("Training Data Shape:", stacked_train.shape)

# Check the shape of the input data for the explainer
print("Explainer Input Shape:", subset_stacked_val_embeddings.shape)
```

Training Data Shape: (28158, 15)

Explainer Input Shape: (100, 5)

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: gru_model = tf.keras.models.load_model("/content/drive/MyDrive/Cyberbullying_
↳Detection/GRU_Model")
```

```
[ ]: cnn_model = tf.keras.models.load_model("/content/drive/MyDrive/Cyberbullying_
↳Detection/CNN_Model")
lstm_model = tf.keras.models.load_model("/content/drive/MyDrive/Cyberbullying_
↳Detection/LSTM_Model")
ensemble_model = tf.keras.models.load_model("/content/drive/MyDrive/
↳Cyberbullying Detection/Stacking_Model.h5")
```

```
[ ]: from lime import lime_text
from lime.lime_text import LimeTextExplainer
from keras.preprocessing.sequence import pad_sequences
```

```
[ ]: def predict_proba(texts):
    encodings = tokenizer(texts, truncation=True, padding='max_length',
↳max_length=100, return_tensors='tf')['input_ids']
    inputs = {'sequences': encodings}
```

```
predictions = gru_model.predict(inputs)
return predictions
```

```
[ ]: label_names = [
    'not bully',
    'troll',
    'sexual',
    'religious',
    'threat'
]
```

```
[ ]: explainer = LimeTextExplainer(class_names = label_names)
```

```
[ ]: for i, sample in enumerate(test_dataset.take(20)):
    input_data, sample_text = sample
    sample_text = sample_text.numpy().decode('utf-8')
    explanation = explainer.explain_instance(sample_text, predict_proba,
    ↪ num_features=10)
    # If you want to see the explanation for each instance, you can visualize
    ↪ it here.
    print(f"Explanation for sample {i+1}:")
    explanation.show_in_notebook()
```

157/157 [=====] - 3s 17ms/step

Explanation for sample 1:

<IPython.core.display.HTML object>

157/157 [=====] - 2s 14ms/step

Explanation for sample 2:

<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step

Explanation for sample 3:

<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step

Explanation for sample 4:

<IPython.core.display.HTML object>

157/157 [=====] - 3s 19ms/step

Explanation for sample 5:

<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step

Explanation for sample 6:

<IPython.core.display.HTML object>

```

157/157 [=====] - 2s 15ms/step
Explanation for sample 7:
<IPython.core.display.HTML object>

157/157 [=====] - 4s 23ms/step
Explanation for sample 8:
<IPython.core.display.HTML object>

157/157 [=====] - 3s 18ms/step
Explanation for sample 9:
<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step
Explanation for sample 10:
<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step
Explanation for sample 11:
<IPython.core.display.HTML object>

157/157 [=====] - 3s 17ms/step
Explanation for sample 12:
<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step
Explanation for sample 13:
<IPython.core.display.HTML object>

157/157 [=====] - 3s 19ms/step
Explanation for sample 14:
<IPython.core.display.HTML object>

157/157 [=====] - 4s 24ms/step
Explanation for sample 15:
<IPython.core.display.HTML object>

157/157 [=====] - 3s 20ms/step
Explanation for sample 16:
<IPython.core.display.HTML object>

157/157 [=====] - 2s 15ms/step
Explanation for sample 17:
<IPython.core.display.HTML object>

157/157 [=====] - 3s 19ms/step
Explanation for sample 18:
<IPython.core.display.HTML object>

```

157/157 [=====] - 2s 15ms/step

Explanation for sample 19:

<IPython.core.display.HTML object>

157/157 [=====] - 3s 16ms/step

Explanation for sample 20:

<IPython.core.display.HTML object>

```
[ ]: def predict_proba(texts):
    encodings = tokenizer(texts, truncation=True, padding='max_length',
    ↪max_length=100, return_tensors='tf')['input_ids']
    inputs = {'sequences': encodings}
    predictions = cnn_model.predict(inputs)
    return predictions

[ ]: for i, sample in enumerate(test_dataset.take(5)):
    input_data, sample_text = sample
    sample_text = sample_text.numpy().decode('utf-8')
    explanation = explainer.explain_instance(sample_text, predict_proba,
    ↪num_features=10)
    # If you want to see the explanation for each instance, you can visualize
    ↪it here.
    print(f"Explanation for sample {i+1}:")
    explanation.show_in_notebook()
```

157/157 [=====] - 4s 3ms/step

Explanation for sample 1:

<IPython.core.display.HTML object>

157/157 [=====] - 1s 3ms/step

Explanation for sample 2:

<IPython.core.display.HTML object>

157/157 [=====] - 1s 3ms/step

Explanation for sample 3:

<IPython.core.display.HTML object>

157/157 [=====] - 1s 4ms/step

Explanation for sample 4:

<IPython.core.display.HTML object>

157/157 [=====] - 0s 3ms/step

Explanation for sample 5:

<IPython.core.display.HTML object>

```
[ ]: def predict_proba(texts):
    encodings = tokenizer(texts, truncation=True, padding='max_length',
    ↪max_length=100, return_tensors='tf')['input_ids']
    inputs = {'sequences': encodings}
    predictions = lstm_model.predict(inputs)
    return predictions

[ ]: for i, sample in enumerate(test_dataset.take(5)):
    input_data, sample_text = sample
    sample_text = sample_text.numpy().decode('utf-8')
    explanation = explainer.explain_instance(sample_text, predict_proba,
    ↪num_features=10)
    # If you want to see the explanation for each instance, you can visualize
    ↪it here.
    print(f"Explanation for sample {i+1}:")
    explanation.show_in_notebook()

[ ]: def predict_proba(texts):
    encodings = tokenizer(texts, truncation=True, padding='max_length',
    ↪max_length=100, return_tensors='tf')['input_ids']
    inputs = {'sequences': encodings}
    predictions = ensemble_model.predict(inputs)
    return predictions

[ ]: for i, sample in enumerate(test_dataset.take(5)):
    input_data, sample_text = sample
    sample_text = sample_text.numpy().decode('utf-8')
    explanation = explainer.explain_instance(sample_text, predict_proba,
    ↪num_features=10)
    # If you want to see the explanation for each instance, you can visualize
    ↪it here.
    print(f"Explanation for sample {i+1}:")
    explanation.show_in_notebook()
```

```
157/157 [=====] - 1s 3ms/step
Explanation for sample 1:
<IPython.core.display.HTML object>

157/157 [=====] - 1s 5ms/step
Explanation for sample 2:
<IPython.core.display.HTML object>

157/157 [=====] - 1s 3ms/step
Explanation for sample 3:
<IPython.core.display.HTML object>
```

157/157 [=====] - 1s 3ms/step

Explanation for sample 4:

<IPython.core.display.HTML object>

157/157 [=====] - 1s 3ms/step

Explanation for sample 5:

<IPython.core.display.HTML object>

Integrated Gradients

```
[ ]: !pip install shap
```

Requirement already satisfied: shap in /usr/local/lib/python3.10/dist-packages (0.44.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.23.5)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.11.4)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.2.2)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (1.5.3)

Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.66.1)

Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (23.2)

Requirement already satisfied: slicer==0.0.7 in /usr/local/lib/python3.10/dist-packages (from shap) (0.0.7)

Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.58.1)

Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.1)

Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap) (0.41.1)

Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2023.3.post1)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (1.3.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.2.0)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->shap) (1.16.0)

```
[ ]: # Assuming you have a specific index for the sequence you want to visualize
sequence_index = 3 # Change this to the index you want to visualize
```

```

# Extract the input data from the test dataset
for i, sample in enumerate(test_dataset):
    if i == sequence_index:
        input_data, _ = sample
        break

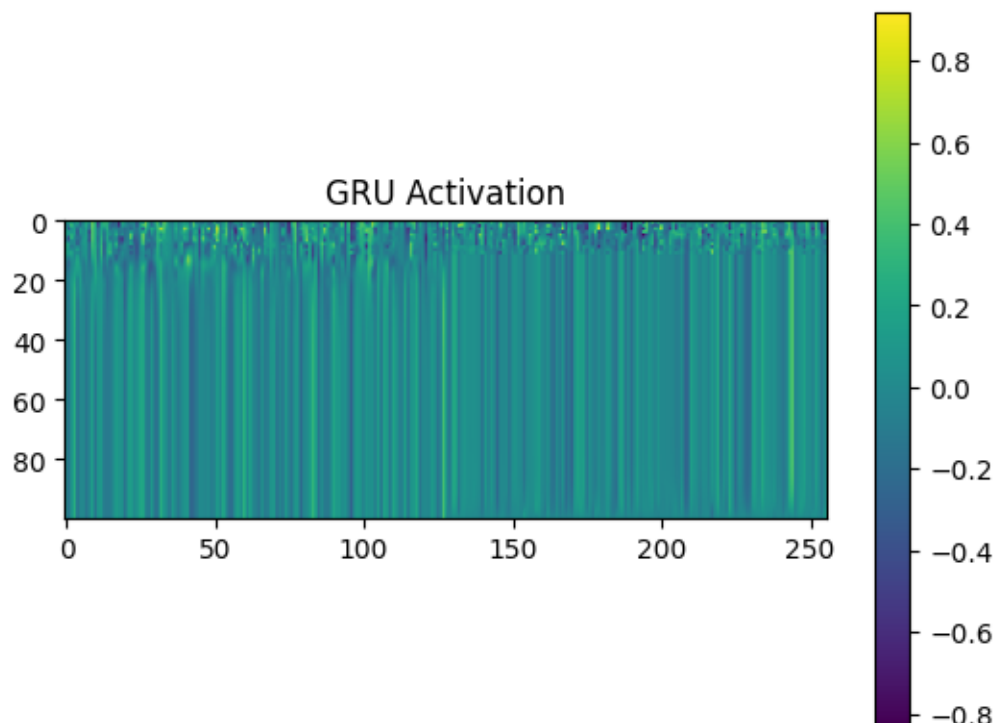
# Ensure the input data is in the correct shape
if len(input_data['input_ids'].shape) == 1:
    sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
else:
    sequences_input = input_data['input_ids']

# Predict with the intermediate model
gru_output = intermediate_layer_model.predict({'sequences': sequences_input})

# Visualize the GRU activation
plt.imshow(gru_output[0])
plt.colorbar()
plt.title('GRU Activation')
plt.show()

```

1/1 [=====] - 8s 8s/step



```
[ ]: cnn_output = intermediate_layer_model.predict({'sequences': sequences_input})
      lstm_output = intermediate_layer_model.predict({'sequences': sequences_input})
      ensemble_output = intermediate_layer_model.predict({'sequences': □
      ↪sequences_input})
```

```
[ ]: # Print layers to identify the desired layer's index
      for i, layer in enumerate(gru_model.layers):
          print(i, layer.name, layer.__class__.__name__)

      # Create the intermediate model (adjust the index accordingly)
      desired_layer_index = 3
      intermediate_layer_model = Model(inputs=gru_model.input, outputs=gru_model.
      ↪layers[desired_layer_index].output)

      input_data = {'input_ids': ...}
      if len(input_data['input_ids'].shape) == 1:
          sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
      else:
          sequences_input = input_data['input_ids']

      # Predict with the intermediate model
      gru_output = intermediate_layer_model.predict({'sequences': sequences_input})

      # Visualize
      plt.imshow(gru_output[0])
      plt.colorbar()
      plt.title('GRU Activation')
      plt.show()
```

```
0 sequences InputLayer
1 embedding_1 Embedding
2 bidirectional_4 Bidirectional
3 dropout_43 Dropout
4 bidirectional_5 Bidirectional
5 dropout_44 Dropout
6 bidirectional_6 Bidirectional
7 dropout_45 Dropout
8 gru_8 GRU
9 dropout_46 Dropout
10 dense_2 Dense
11 dropout_47 Dropout
12 dense_3 Dense
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-192-9abc00062de2> in <cell line: 10>()
      8
```



```

    9 input_data = {'input_ids': ...}
---> 10 if len(input_data['input_ids'].shape) == 1:
    11     sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
    12 else:

```

AttributeError: 'ellipsis' object has no attribute 'shape'

```

[ ]: # Print layers to identify the desired layer's index
for i, layer in enumerate(cnn_model.layers):
    print(i, layer.name, layer.__class__.__name__)

# Create the intermediate model (adjust the index accordingly)
desired_layer_index = 3
intermediate_layer_model = Model(inputs=cnn_model.input, outputs=cnn_model.
    ↳ layers[desired_layer_index].output)

input_data = {'input_ids': ...}
if len(input_data['input_ids'].shape) == 1:
    sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
else:
    sequences_input = input_data['input_ids']

# Predict with the intermediate model
cnn_output = intermediate_layer_model.predict({'sequences': sequences_input})

# Visualize
plt.imshow(gru_output[0])
plt.colorbar()
plt.title('CNN Activation')
plt.show()

```

```

[ ]: # Print layers to identify the desired layer's index
for i, layer in enumerate(cnn_model.layers):
    print(i, layer.name, layer.__class__.__name__)

# Create the intermediate model (adjust the index accordingly)
desired_layer_index = 3
intermediate_layer_model = Model(inputs=cnn_model.input, outputs=cnn_model.
    ↳ layers[desired_layer_index].output)

input_data = {'input_ids': ...}
if len(input_data['input_ids'].shape) == 1:
    sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
else:
    sequences_input = input_data['input_ids']

```

```

# Predict with the intermediate model
cnn_output = intermediate_layer_model.predict({'sequences': sequences_input})

# Visualize
plt.imshow(gru_output[0])
plt.colorbar()
plt.title('CNN Activation')
plt.show()


```

WARNING:tensorflow:6 out of the last 790 calls to <function Model.make_predict_function.<locals>.predict_function at 0x79d7b510aa70> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

```

0 sequences InputLayer
1 embedding_4 Embedding
2 conv1d_2 Conv1D
3 max_pooling1d_2 MaxPooling1D
4 conv1d_3 Conv1D
5 max_pooling1d_3 MaxPooling1D
6 flatten_1 Flatten
7 dense_8 Dense
8 dense_9 Dense
9 dense_10 Dense
10 dense_11 Dense
1/1 [=====] - 0s 78ms/step

```



```
[ ]: # Print layers to identify the desired layer's index
for i, layer in enumerate(lstm_model.layers):
    print(i, layer.name, layer.__class__.__name__)

# Create the intermediate model (adjust the index accordingly)
desired_layer_index = 3
intermediate_layer_model = Model(inputs=lstm_model.input, outputs=lstm_model.
    ↳ layers[desired_layer_index].output)

if len(input_data['input_ids'].shape) == 1:
    sequences_input = np.expand_dims(input_data['input_ids'], axis=0)
else:
    sequences_input = input_data['input_ids']


# Predict with the intermediate model
cnn_output = intermediate_layer_model.predict({'sequences': sequences_input})

# Visualize
plt.imshow(gru_output[0])
plt.colorbar()
plt.title('LSTM Activation')
plt.show()
```

```

0 sequences InputLayer
1 embedding_2 Embedding
2 bidirectional_7 Bidirectional
3 dropout_48 Dropout
4 lstm_1 LSTM
5 dropout_49 Dropout
6 dense_4 Dense
7 dropout_50 Dropout
8 dense_5 Dense
1/1 [=====] - 2s 2s/step

```



```

[ ]: # Print layers to identify the desired layer's index
for i, layer in enumerate(ensemble_model.layers):
    print(i, layer.name, layer.__class__.__name__)

desired_layer_index = 3
intermediate_layer_model = Model(inputs=ensemble_model.input,
    ↪ outputs=ensemble_model.layers[desired_layer_index].output)

if len(input_data['input_ids'].shape) == 1:
    sequences_input = np.expand_dims(input_data['input_ids'], axis=0)

```

```

else:
    sequences_input = input_data['input_ids']

    # Predict with the intermediate model
    cnn_output = intermediate_layer_model.predict({'sequences': sequences_input})


    # Visualize
    plt.imshow(gru_output[0])
    plt.colorbar()
    plt.title('Ensemble Activation')
    plt.show()

```

```

0 sequences InputLayer
1 embedding_5 Embedding
2 conv1d_2 Conv1D
3 max_pooling1d_2 MaxPooling1D
4 conv1d_3 Conv1D
5 max_pooling1d_3 MaxPooling1D
6 flatten_1 Flatten
7 dense_12 Dense
8 dense_13 Dense
9 dense_14 Dense
10 dense_15 Dense
1/1 [=====] - 0s 48ms/step

```



[]: